

# Local Area Networks: Ethernet, Switching


Note: The slides are adapted from the materials from Prof. Richard Han at CU Boulder and Profs. Jennifer Rexford and Mike Freedman at Princeton University, and the networking book (Computer Networking: A Top Down Approach) from Kurose and Ross.

# Ethernet and Switching


# Goals of Today's Lecture

- Devices that shuttle data at different layers
  - Repeaters and hubs
  - Bridges and switches
  - Routers
- Switch protocols and mechanisms
  - Dedicated access and full-duplex transfers
  - Cut-through switching
  - Self learning of the switch table
  - Spanning trees
- Virtual LANs (VLANs)

# Fully-connected links



# Shared broadcast medium




It's all about resource allocation

# Three Ways to Share the Media

- Channel partitioning MAC protocols:
  - Share channel efficiently and fairly at high load
  - Inefficient at low load: unused go idle
- “Taking turns” protocols
  - Eliminates empty slots without causing collisions
  - Vulnerable to failures
- Random access MAC protocols
  - Efficient at low load: single node can fully utilize channel
  - High load: collision overhead


# Hubs:

## Joining broadcast mediums



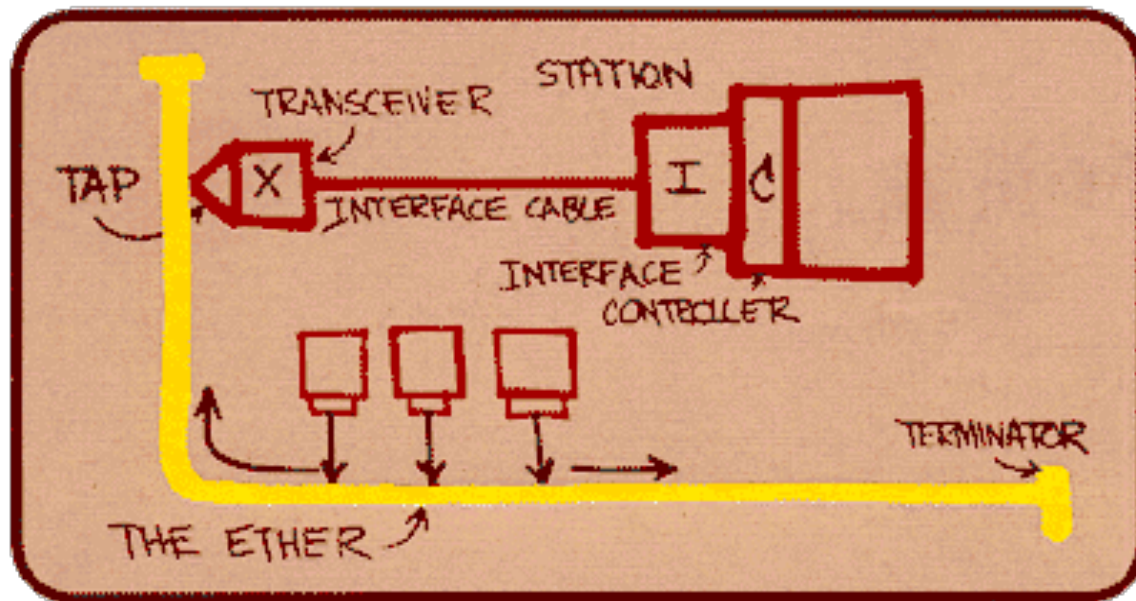


# Bridges / Switches: Isolating broadcast mediums




# Ethernet

- Dominant wired LAN technology, first widely used
- Simpler, cheaper than token LANs and ATM
- Kept up with speed race: 10 Mbps – 10 Gbps



Metcalfe's  
Ethernet  
sketch

# Ethernet Frame Structure

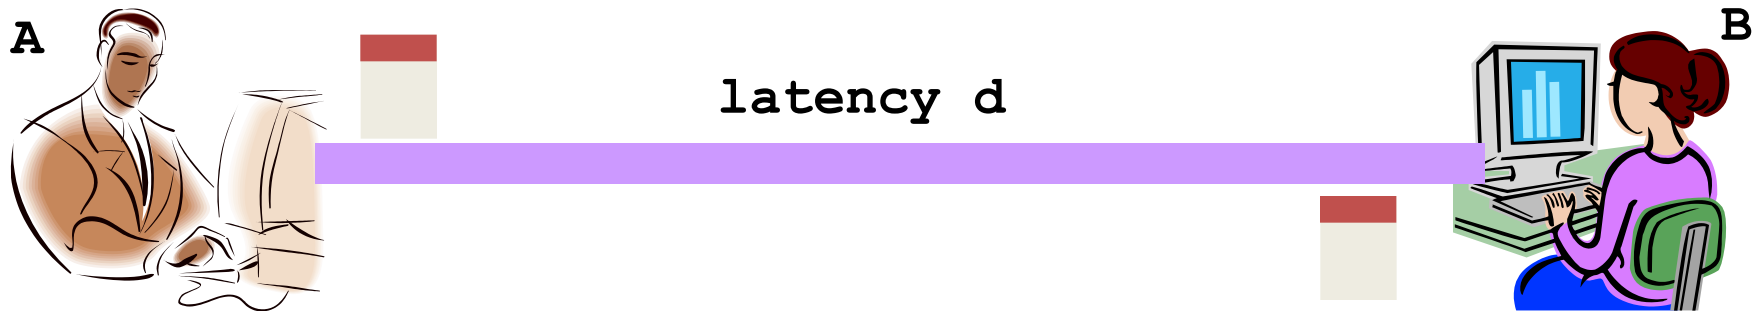


- **Preamble:** synchronization:  $(10101010)^7 10101011$
- **Addresses:** 6-byte source and dest MAC addresses
  - Adaptor passes frame to OS stack if destination matches adaptor or is broadcast address; otherwise, discard frame
- **Type:** higher-layer protocol (IP, AppleTalk, ...)
- **Error detection:** CRC: cyclic redundancy check
- **Best effort:** Connectionless, unreliable

# Ethernet Uses CSMA/CD


- Carrier Sense: wait for link to be idle before transmit
- Collision Detection: listen while transmitting
  - No collision: transmission complete
  - Collision: abort and send jam signal
- Random access: exponential back-off
  - After collision, wait a random time before retry
  - After  $m^{\text{th}}$  collision, choose  $K$  randomly from  $\{0, \dots, 2^m - 1\}$
  - ... and wait for  $K * 64$  byte times before retry

# Limitations on Ethernet Length



- Latency depends on physical length of link
  - Time to propagate a packet from one end to the other
- Suppose A sends a packet at time  $t$ 
  - And B sees an idle line just before time  $t+d$ , so transmits
- B detects a collision, and sends jamming signal
  - But A doesn't see collision till  $t+2d$

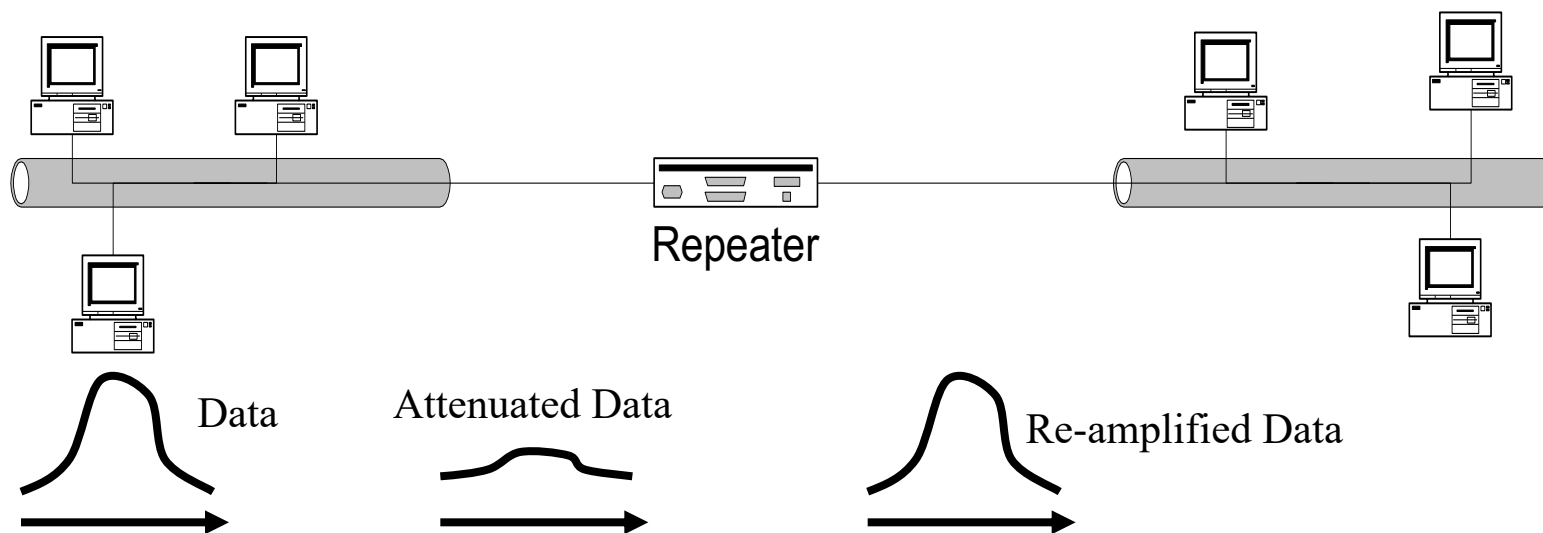
# Limitations on Ethernet Length



- **A needs to wait for time  $2d$  to detect collision**
  - So, A should keep transmitting during this period
  - ... and keep an eye out for a possible collision
- **Imposes restrictions on Ethernet**
  - Max length of wire: 2500 meters
  - Min length of packet: 512 bits (64 bytes)


# Physical Layer: Repeaters

- Distance limitation in local-area networks
  - Electrical signal becomes weaker as it travels
  - Imposes a limit on the length of a LAN
- Repeaters join LANs together
  - Analog electronic device
  - Monitors signals on each LAN and transmits amplified copies



# Physical Layer: Hubs

- **Joins multiple input lines electrically**
  - Designed to hold multiple line cards
  - Do not necessarily amplify the signal
  - Monitoring and fault isolation are supported
- **Very similar to repeaters**
  - Also operates at the physical layer






# An Actual Ethernet Hub

8 ports: connect hosts here

Uplink: connect to  
another hub





# Limitations of Repeaters and Hubs

- One large shared link
  - Each bit sent everywhere, aggregate throughput limited
- Cannot support multiple LAN technologies
  - Does not buffer or interpret frames
  - So, can't interconnect different rates or formats
- Limitations on maximum nodes and distances

# Switching for resource isolation


# Link Layer: Bridges and Switches

- Connects two or more LANs at the link layer
  - Extracts destination address from the frame
  - Looks up the destination in a table, forwards to appropriate
- Each segment can carry its own traffic
  - Concurrent traffic between LANs/host: A to B while D to C
- Bridge: connecting LANs; Switches: connecting hosts



# Bridges/Switches: Traffic Isolation

- Switch breaks subnet into LAN segments
- Switch filters packets
  - Frame only forwarded to the necessary segments
  - Segments can support separate transmissions



# High-density switching



rack



48-port switch



Facebook rack

- Each rack has 42 U (“pizza boxes”)
- Typically servers + 1-2 “top-of-rack” switch(es)

# Advantages Over Hubs/Repeaters

- Only forwards frames as needed
  - E.g. to destination segments or for broadcast traffic
  - Reduces unnecessary traffic on segments
- Extends the geographic span of the network
  - Ethernet collisions (and distance limitations) only on segment
- Improves privacy by limiting scope of frames
  - Hosts can only “snoop” the traffic traversing *their* segment
- Can join segments using different technologies

# Disadvantages Over Hubs/Repeaters

- Delay in forwarding frames
  - Bridge/switch must receive frame, parse, lookup, and send
  - Storing and forwarding the packet introduces delay
  - Sol'n: **cut-through switching** (start send after receive header)
- Need to learn where to forward frames
  - Forwarding table: destination MAC → outgoing interface
  - Needs to construct forwarding table, ideally w/o static config
  - Sol'n: **self-learning**
- Higher cost
  - More complicated devices that cost more money




# Switches vs. Hubs

- Compared to hubs, Ethernet switches support
  - (a) Larger geographic span
  - (b) Similar span
  - (c) Smaller span
- Compared to hubs, switches provides
  - (a) Higher load on links
  - (b) Less privacy
  - (c) Heterogeneous communication technologies


# Self Learning: Building the Table

- When a frame arrives
  - Inspect *source* MAC address
  - Associate addr with *incoming* interface/port
  - Store mapping in forwarding table
  - Use TTL field to eventually forget mapping



# Self Learning: Handling Misses

- When frame arrives with unfamiliar destination
  - Forward frame out all interfaces except source
  - Hopefully, won't happen very often



# Switch Filtering/Forwarding

When switch receives a frame:

index switch table using MAC dest address

if ( entry found for destination ) then

if ( dest on segment from which frame arrived ) then

drop the frame

else

forward the frame on interface indicated


else flood



forward on all but the interface  
on which the frame arrived


# Table for Forwarding and Filtering

- Label the interfaces into/out of the bridge
- When a frame arrives, store the **source address** and **the originating interface** from which the frame came



Frame to Z →	Src Addr	Originating Interface	Time
	U	1	--


# Frame Forwarding Rules



Event	Originating Interface 1	Originating Interface 2	Bridge Forwarding Action
Boot state	--	--	
U sends to V	U	--	Forward to all outgoing (e.g. 2)
V sends to U	U,V	--	none
Z sends to U	U,V	Z	Forward to 1
Z sends to Y	U,V	Z	Forward to all outgoing (e.g. 1)
Y sends to V	U,V	Y,Z	Forward to 1


# Flooding Can Lead to Loops

- E.g., if the network contains a cycle of switches
- Either accidentally or by design for higher reliability



- **Solution: Spanning Tree**
  - Ensure the topology has no loops
  - Avoid using some of the links when flooding
  - Spanning tree: Sub-graph that covers all vertices but contains no cycles

# Spanning Trees




- **Solution: Spanning Tree**
  - Ensure the topology has no loops
  - Avoid using some of the links when flooding
  - Spanning tree: Sub-graph that covers all vertices but contains no cycles



# Constructing a Spanning Tree

- **Distributed algorithm**
  - Switches cooperate to build, auto-adapt on failures
- **Key ingredients of the algorithm**
  - Switches elect a “root” (e.g. one w/ smallest ID)
  - Each determines if interface is on shortest path from root, excludes if not
  - Learned via messages from peers
    - (root Y, distance d, from X)
  - Reacts to root/switch/link failures
    - Path entries have TTL (i.e. soft state)
    - Root periodically reannounces




# Steps in Spanning Tree Algorithm

- **Initially, each switch thinks it is the root**
  - Switch sends a message out every interface
  - ... identifying itself as the root with distance 0
  - Example: switch X announces (X, 0, X)
- **Switches update their view of the root**
  - Upon receiving a message, check the root id
  - If the new id is smaller, start viewing that switch as root
- **Switches compute their distance from the root**
  - Add 1 to the distance received from a neighbor
  - Identify interfaces not on a shortest path to the root
  - ... and exclude them from the spanning tree


# Example From Switch #4' s Viewpoint

- **Switch #4 thinks it is the root**
  - Sends (4, 0, 4) message to 2 and 7
- **Then, switch #4 hears from #2**
  - Receives (2, 0, 2) message from 2
  - ... and thinks that #2 is the root
  - And realizes it is just one hop away
- **Then, switch #4 hears from #7**
  - Receives (2, 1, 7) from 7
  - And realizes this is a longer path
  - So, prefers its own one-hop path
  - And removes 4-7 link from the tree




# Example From Switch #4's Viewpoint

- **Switch #2 hears about switch #1**
  - Switch 2 hears (1, 1, 3) from 3
  - Switch 2 starts treating 1 as root
  - And sends (1, 2, 2) to neighbors
- **Switch #4 hears from switch #2**
  - Switch 4 starts treating 1 as root
  - And sends (1, 3, 4) to neighbors
- **Switch #4 hears from switch #7**
  - Switch 4 receives (1, 3, 7) from 7
  - And realizes this is a longer path
  - So, prefers its own three-hop path
  - And removes 4-7 link from the tree



# STP Problem



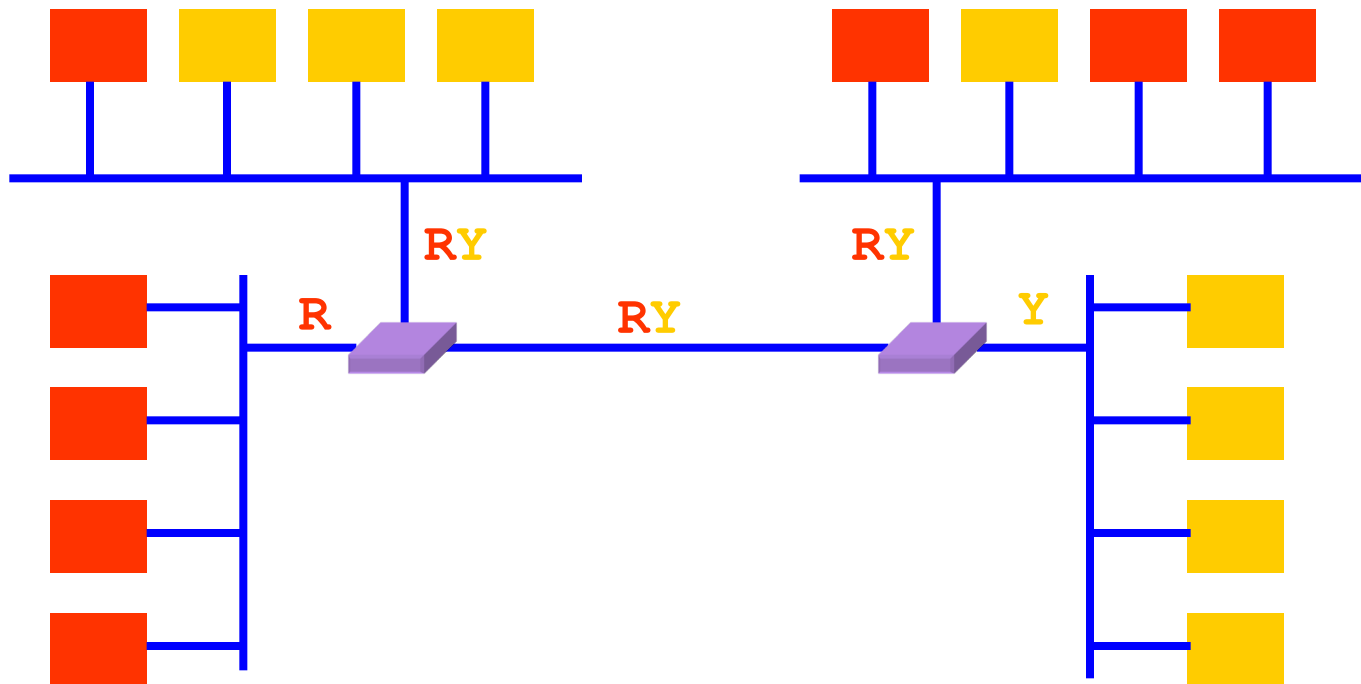
# Evolution Toward Virtual LANs

- In the olden days...
  - Thick cables snaked through cable ducts in buildings
  - Every computer was plugged in
  - All people in adjacent offices were on same LAN
- More recently due to hubs and switches...
  - Every office connected to central wiring closets
  - Flexibility in mapping offices to different LANs
- Evolution to grouping users based on org structure, not physical layout of building

# Why Group by Org Structure?

- **Security**
  - Ethernet is a shared media
  - Interfaces can be put in “promiscuous” mode to see all traffic
- **Load**
  - Some LAN segments are more heavily used than others
    - E.g., researchers can saturate own segment, but not others
  - May be natural locality of communication
    - E.g., traffic between people in the same research group
- **But people move, organizations changes**
  - Physical rewiring is a huge pain!

# Virtual LANs




**Red VLAN and Yellow VLAN**

**Switches forward traffic as needed**



# Virtual LANs



**Red VLAN and Yellow VLAN**

**Switches forward traffic as needed**

# Making VLANs Work

- Switches need configuration tables
  - Saying which VLANs are accessible via which interfaces
- Approaches to mapping to VLANs
  - VLAN color per interface
    - Only if all hosts on segment belong to same VLAN
  - VLAN color per MAC address
- Changing the Ethernet header
  - Adding a field for a VLAN tag
  - VLAN tag added/removed by switches
    - Hosts unaware (backwards compat), cannot spoof (security)

# Comparing Hubs, Switches, Routers

	Hub / Repeater	Bridge / Switch	IP Router
Traffic isolation	No	Yes	Yes
Plug and Play	Yes	Yes	No
Efficient routing	No	No	Yes