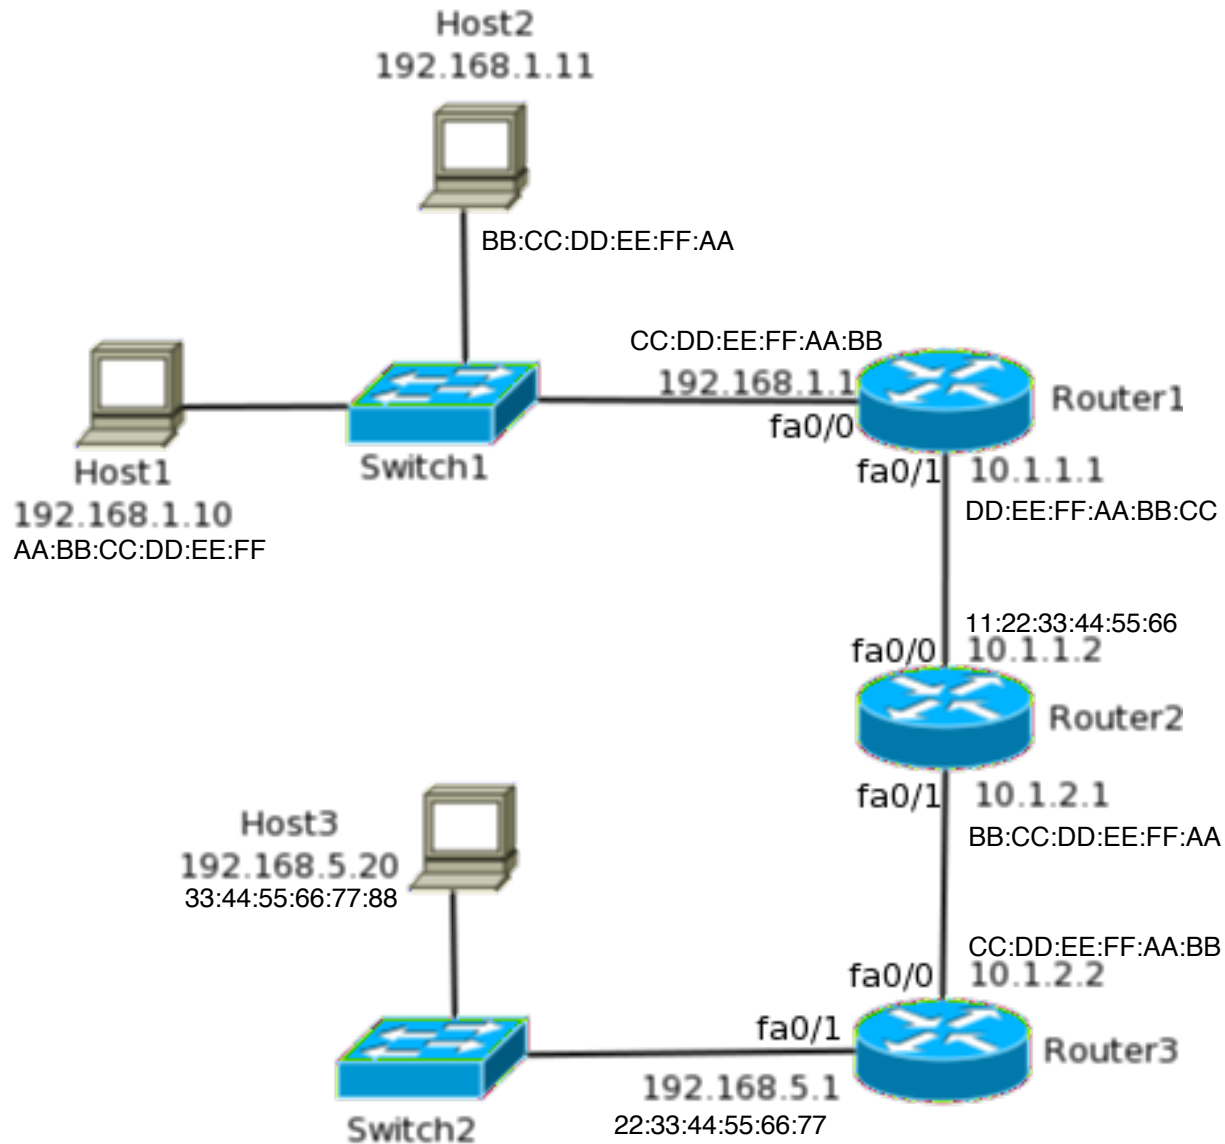


Routing

Note: The slides are adapted from the materials from Prof. Richard Han at CU Boulder and Profs. Jennifer Rexford and Mike Freedman at Princeton University, and the networking book (Computer Networking: A Top Down Approach) from Kurose and Ross.

Review

- DNS
- ARP
- DHCP



Goals of Today's Lecture


- Inside a router
 - Control plane: routing protocols
 - Data plane: packet forwarding
- Path selection
 - Minimum-hop/cost and shortest-path routing
 - Algorithms: Link-state vs. Distance vector routing
- Topology change
 - Using beacons to detect topology changes
 - Propagating topology information

What is Routing?

- A famous quotation from RFC 791

“A *name* indicates what we seek.
An *address* indicates where it is.
A *route* indicates how we get there.”

-- Jon Postel

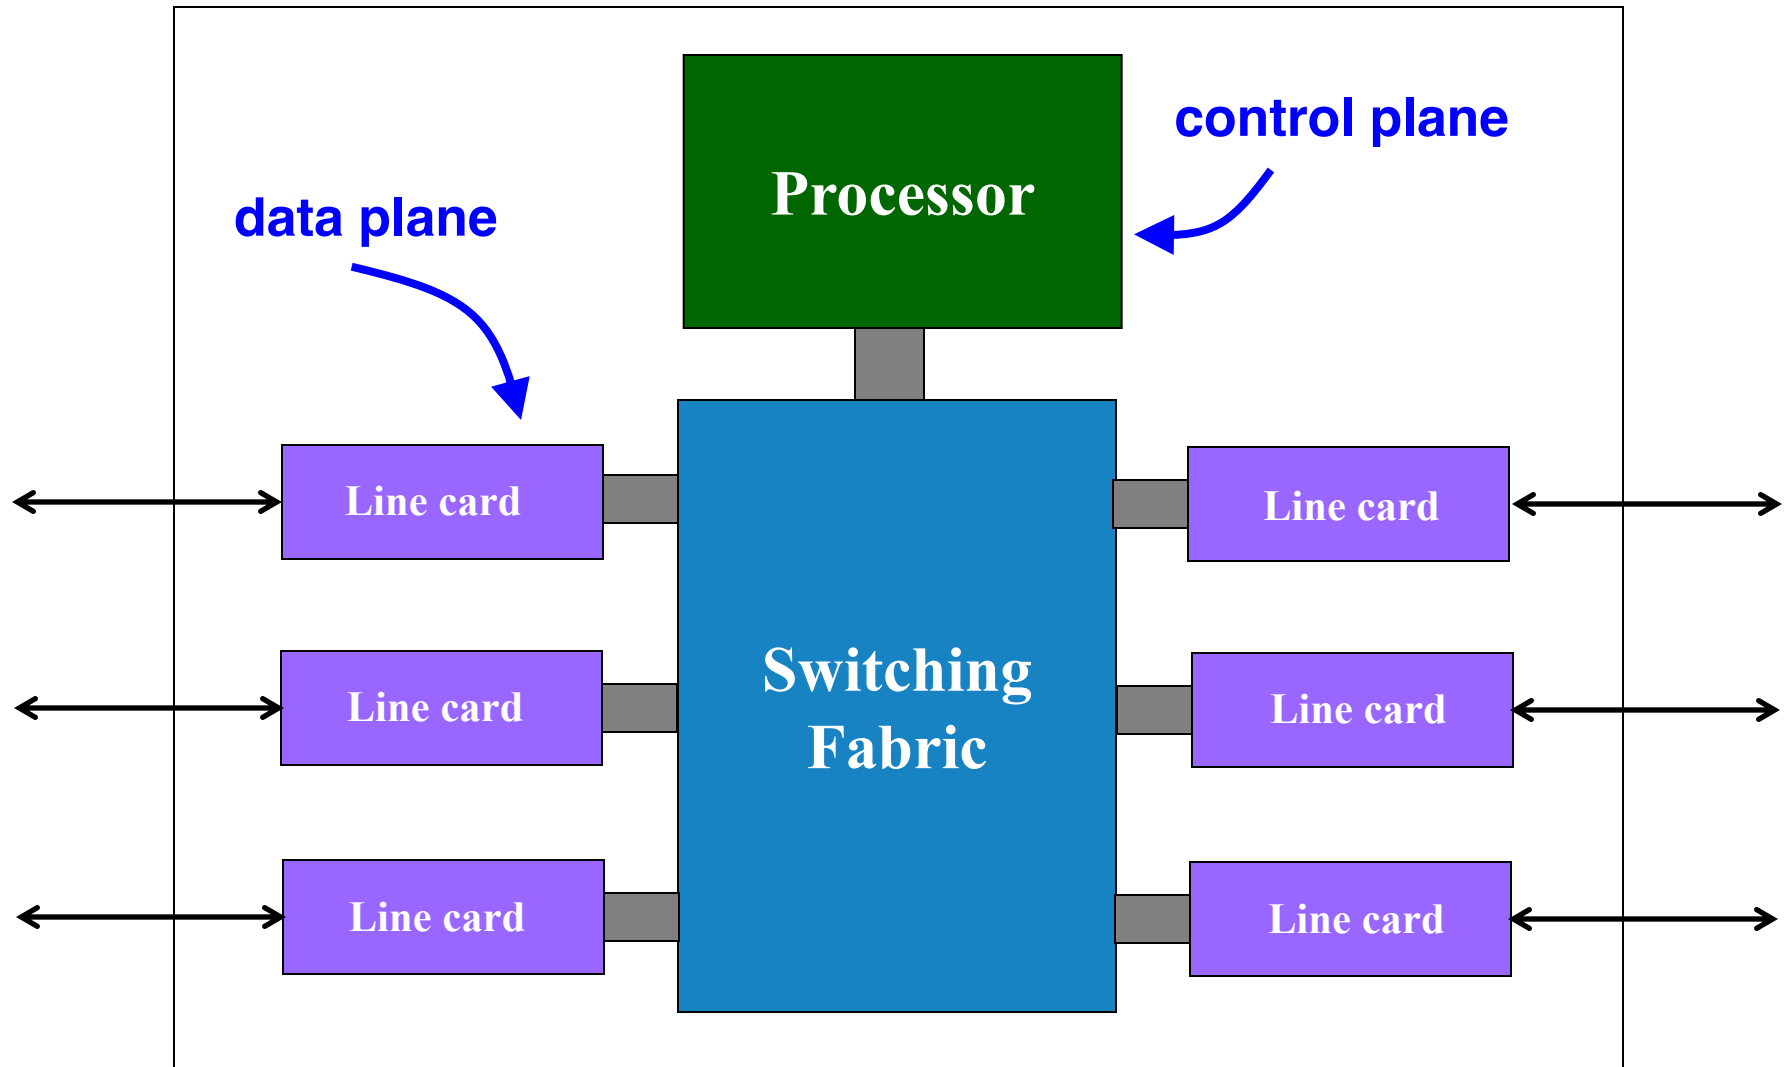


Routing vs. Forwarding

- **Routing:** control plane
 - Computing paths the packets will follow
 - Routers talking amongst themselves
 - Individual router *creating* a forwarding table
- **Forwarding:** data plane
 - Directing a data packet to an outgoing link
 - Individual router *using* a forwarding table



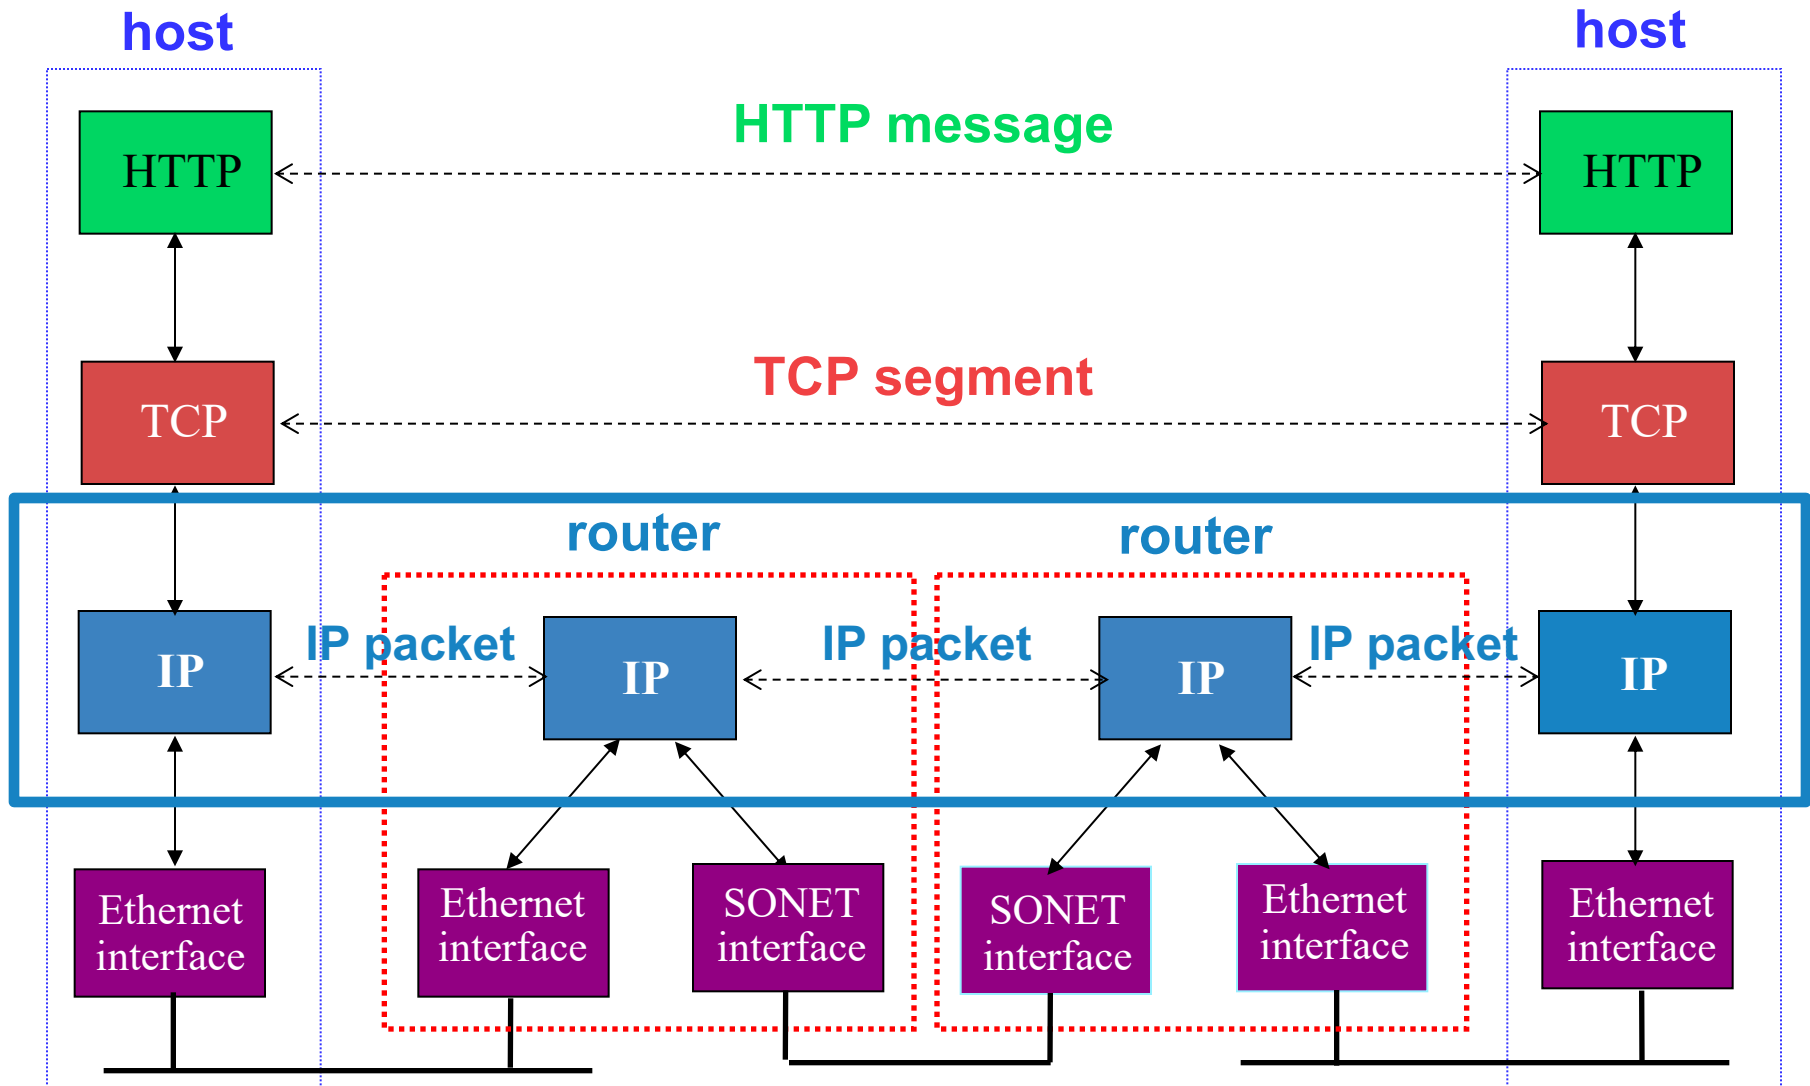
Data and Control Planes



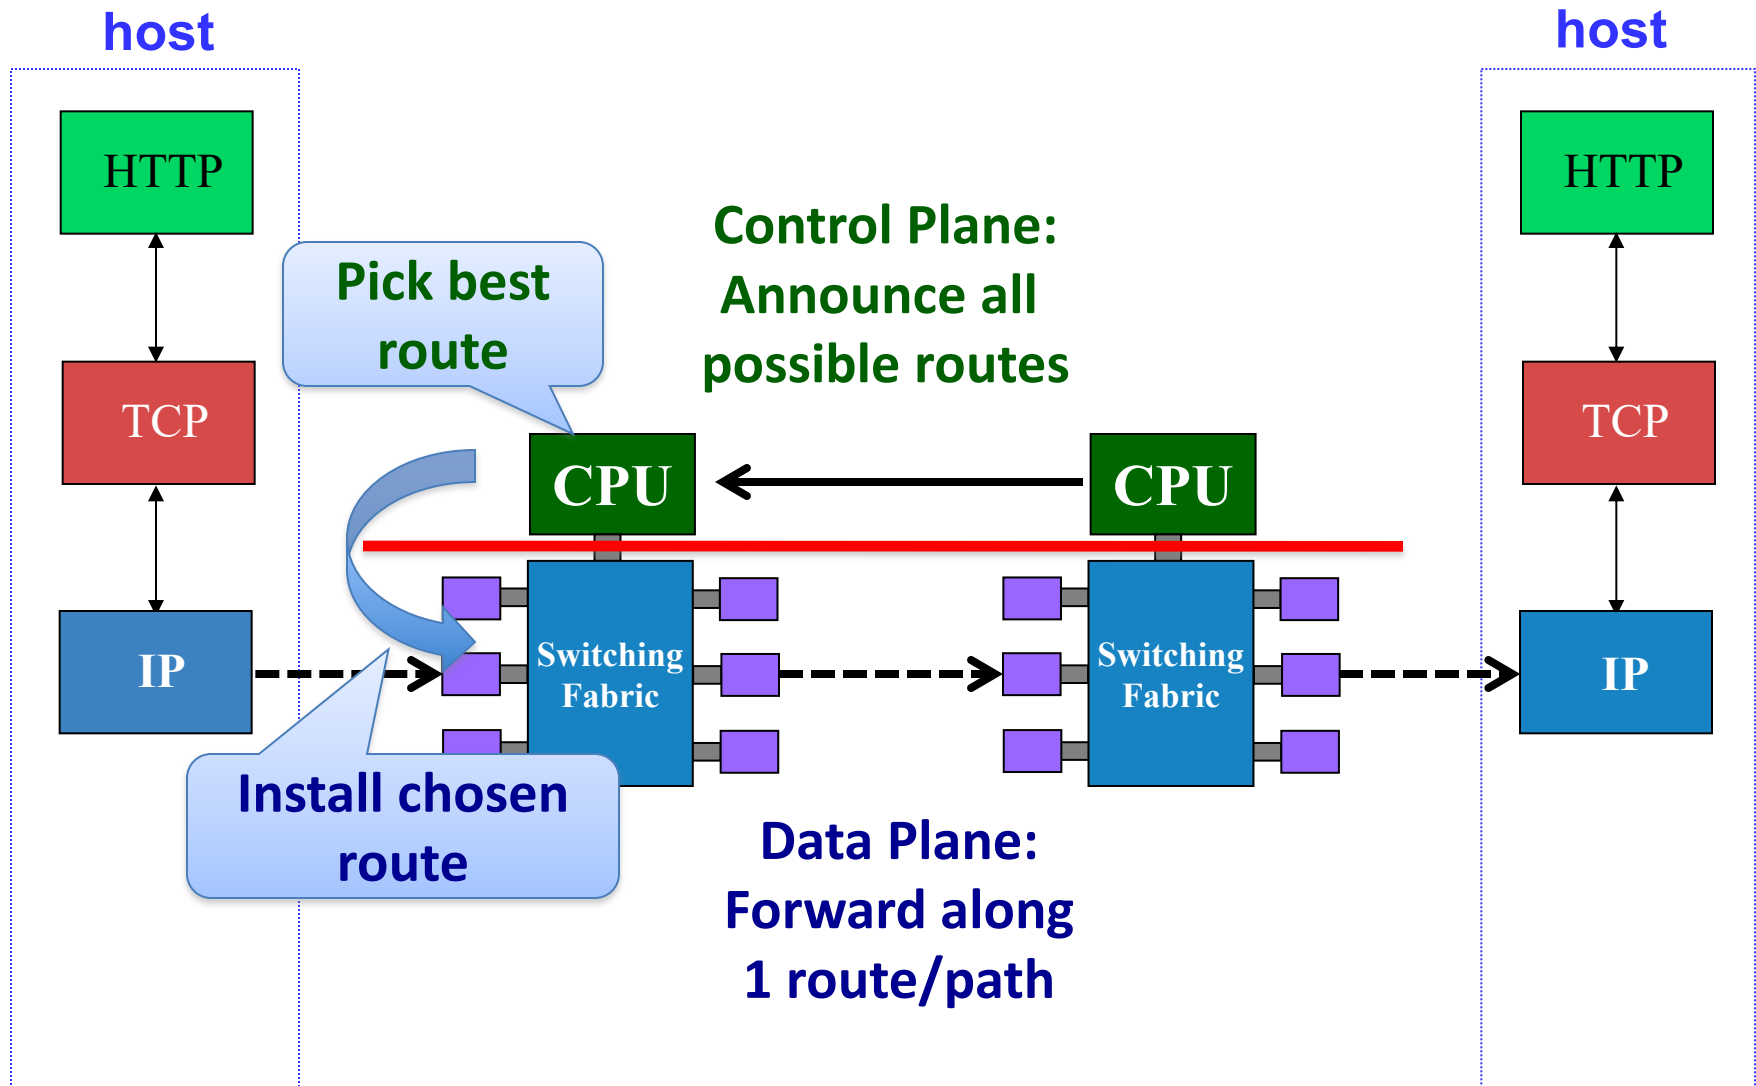
Where do Forwarding Tables Come From?

- Routers have forwarding tables
 - Map IP prefix to outgoing link(s)
- Entries can be statically configured
 - E.g., “map 12.34.158.0/24 to Serial0/0.1”
- But, this doesn't adapt
 - To failures
 - To new equipment
 - To the need to balance load
- That is where routing protocols come in

Recall the Internet layering model




Recall the Internet layering model



Computing Paths Between Routers

- Routers need to know two things
 - Which router to use to reach a destination prefix
 - Which outgoing interface to use to reach that router




- Today's class: how routers reach each other
 - How *u* knows how to forward packets toward *z*

Computing the Shortest Paths

Assuming you already know
the topology


Shortest-Path Routing

- Path-selection model
 - Destination-based
 - Load-insensitive (e.g., static link weights)
 - Minimum hop count or sum of link weights



Shortest-Path Problem

- **Given: network topology with link costs**
 - $c(x,y)$: link cost from node x to node y
 - Infinity if x and y are not direct neighbors
- **Compute: least-cost paths to all nodes**
 - From a given source u to all other nodes
 - $p(v)$: predecessor node along path from source to v



Link-State Routing

Link State: Dijkstra's Algorithm

- Flood the topology information to all nodes
- Each node computes shortest paths to other nodes

Initialization





```
S = {u}
for all nodes v
  if (v is adjacent to u)
    D(v) = c(u,v)
  else D(v) =  $\infty$ 
```


Loop

```
add w with smallest D(w) to S
update D(v) for all adjacent v:
  D(v) = min{D(v), D(w) + c(w,v)}
until all nodes are in S
```

Used in OSPF and IS-IS


Link-State Routing Example






Link State: Shortest-Path Tree

- Shortest-path tree from u
- Forwarding table at u



	link
v	(u,v)
w	(u,w)
x	(u,w)
y	(u,v)
z	(u,v)
s	(u,w)
t	(u,w)


Link State: Shortest-Path Tree




Find shortest path t to v

- Forwarding table entry at t
(A) (t,x) (B) (t, s)
- Distance from t to v
(A) 6 (B) 7 (C) 8 (D) 9

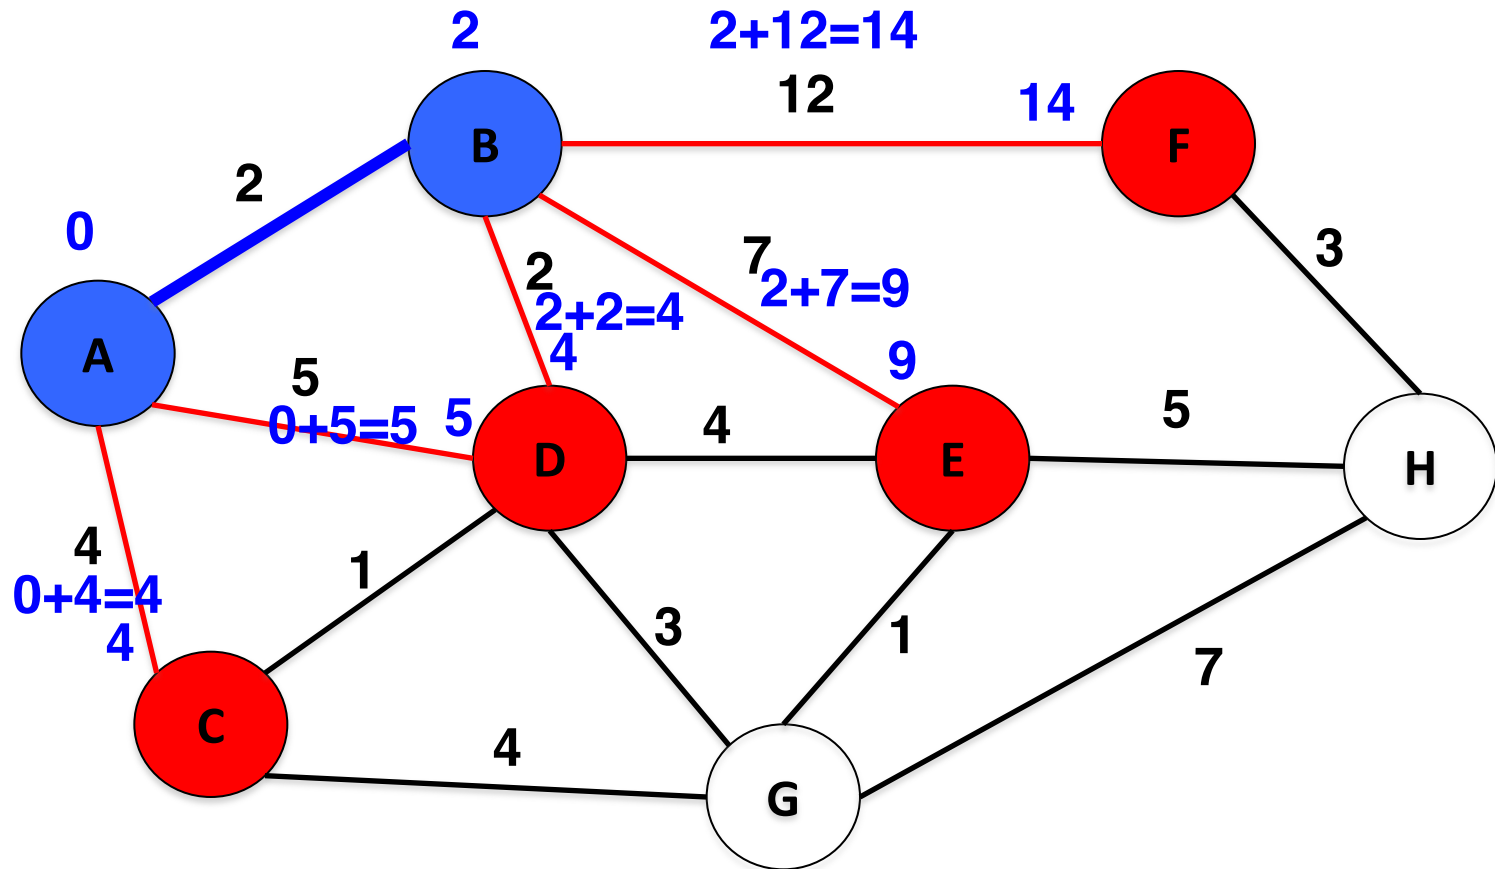
Link-State Algorithm Example 2




Link-State Algorithm Example 2 (Cont.)




Link-State Algorithm Example 2 (Cont.)



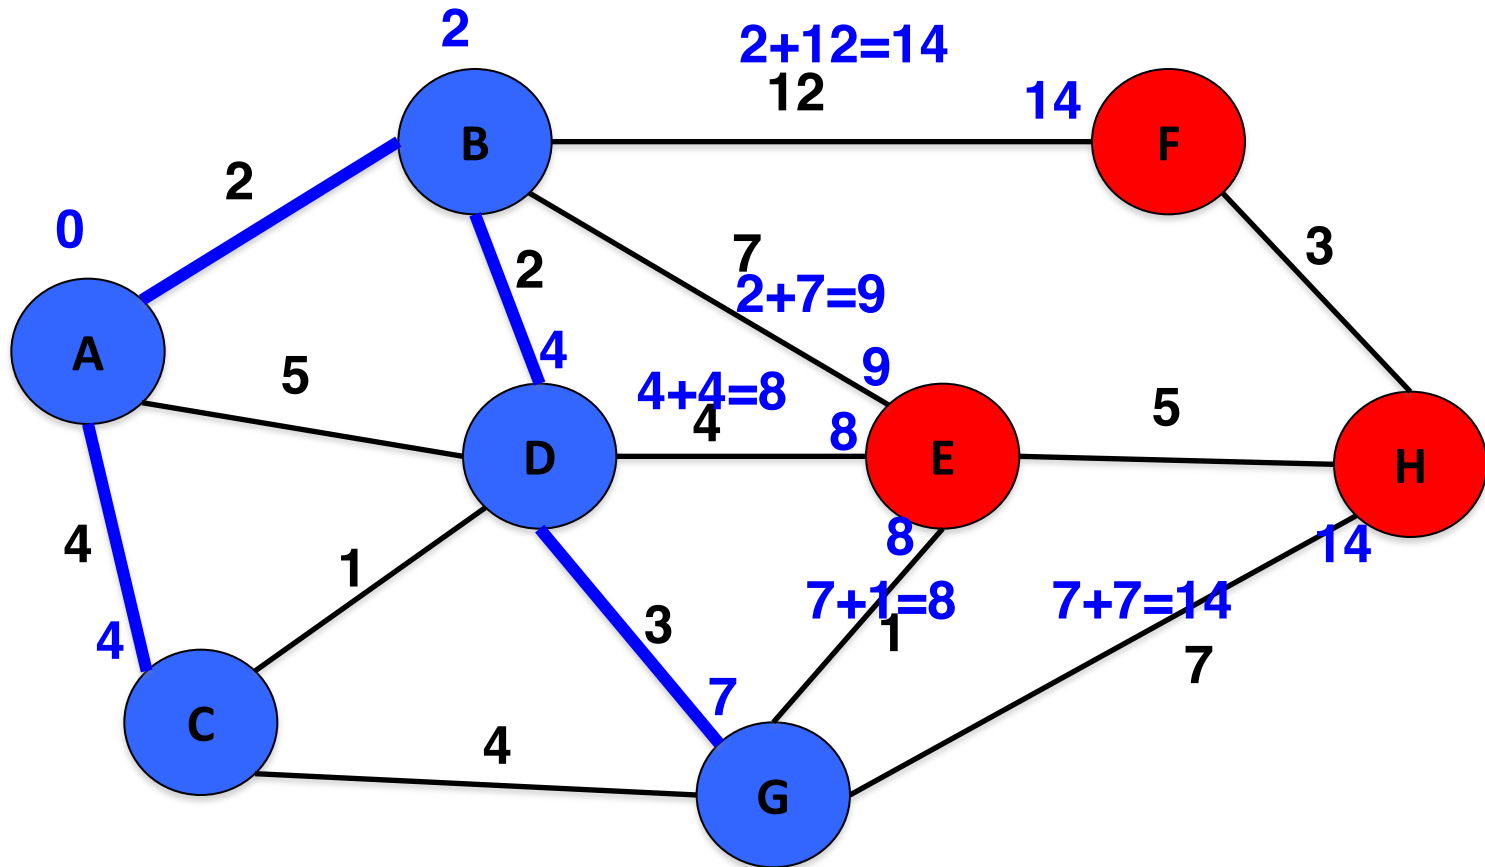
Link-State Algorithm Example 2 (Cont.)



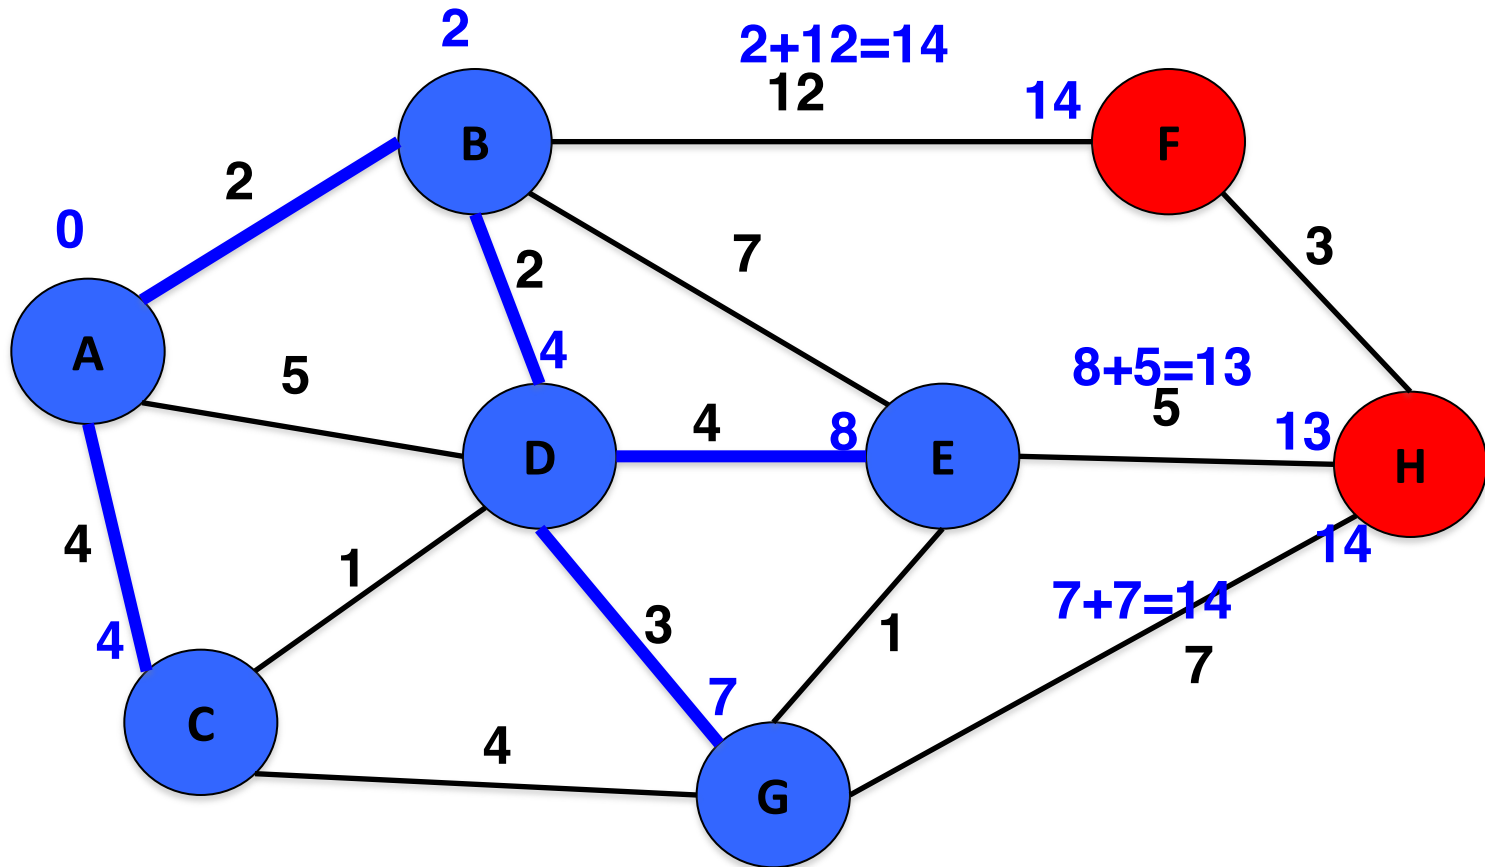
Link-State Algorithm Example 2 (Cont.)



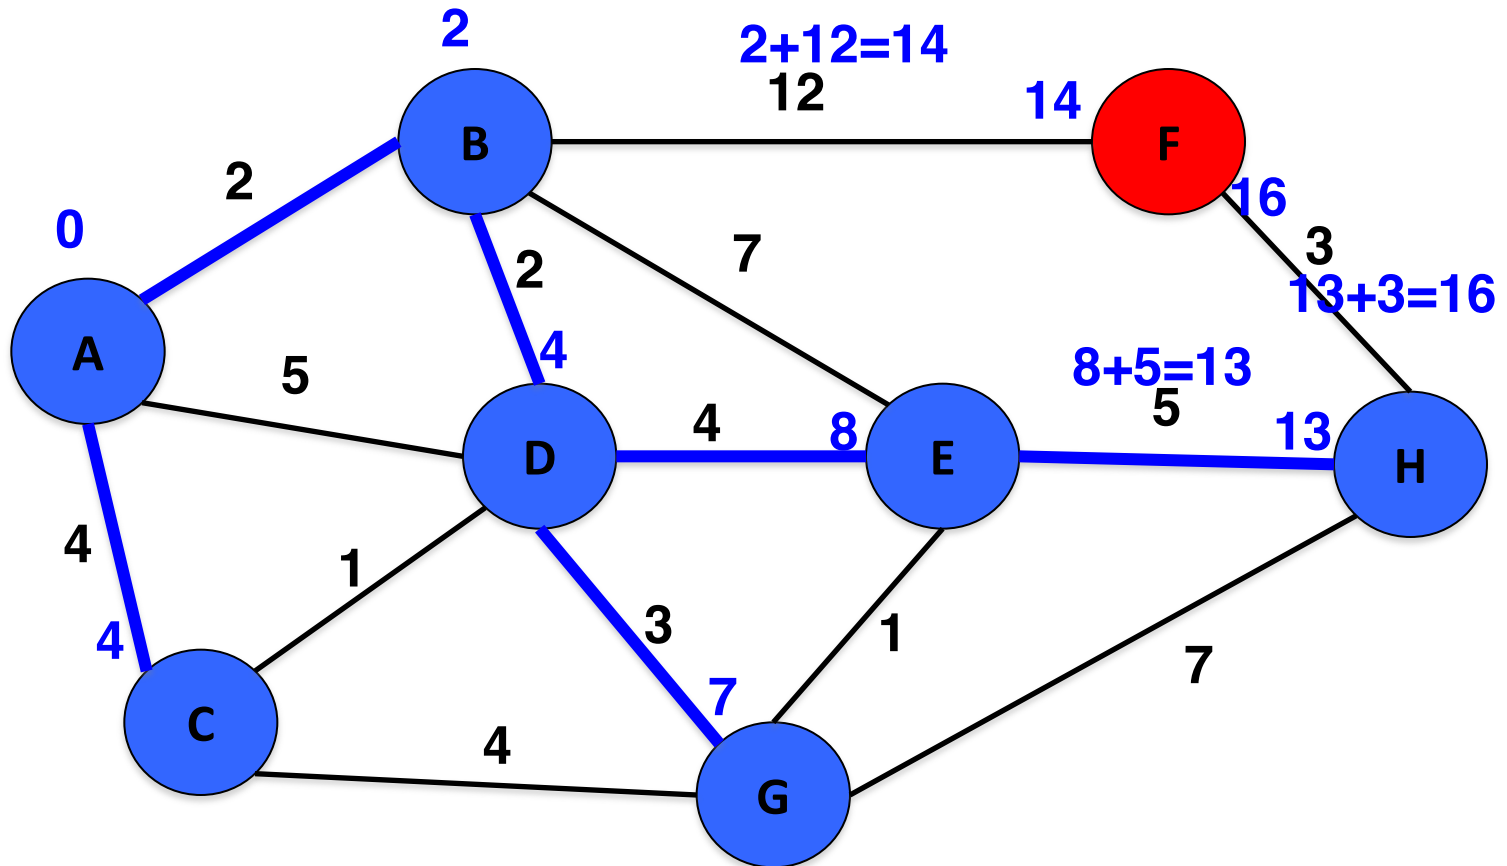
Link-State Algorithm Example 2 (Cont.)




Link-State Algorithm Example 2 (Cont.)




Link-State Algorithm Example 2 (Cont.)



Link-State Algorithm Example 2 (Cont.)




Link-State Algorithm Example 2 (Cont.)



Distance Vector Routing

Distance Vector: Bellman-Ford Algo


- Define distances at each node x
 - $d_x(y)$ = cost of least-cost path from x to y
- Update distances based on neighbors
 - $d_x(y) = \min \{c(x,v) + d_v(y)\}$ over all neighbors v



$$d_u(z) = \min \{ c(u,v) + d_v(z), \\ c(u,w) + d_w(z) \}$$


Used in RIP and EIGRP

Distance Vector Example



$$d_y(z) = 1$$


$$d_x(z) = 4$$




$$d_v(z) = \min \{ 2 + d_y(z), 1 + d_x(z) \}$$

$$= 3$$

Distance Vector Example (Cont.)




$$d_w(z) = \min\{ 1+d_x(z), \\ 4+d_s(z), \\ 2+d_u(z) \} \\ = 5$$




$$d_u(z) =$$

(A) 5 (B) 6 (C) 7

Distance Vector Example (Cont.)



$$d_w(z) = \min \{ 1+d_x(z), \\ 4+d_s(z), \\ 2+d_u(z) \} \\ = 5$$



$$d_u(z) = \min \{ 3+d_v(z), \\ 2+d_w(z) \} \\ = 6$$

Distance Vector Example 2: Step 1

Optimum 1-hop paths

Table for A			Table for B		
Dst	Cst	Hop	Dst	Cst	Hop
A	0	A	A	4	A
B	4	B	B	0	B
C	∞	—	C	∞	—
D	∞	—	D	3	D
E	2	E	E	∞	—
F	6	F	F	1	F





Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	∞	—	A	∞	—	A	2	A	A	6	A
B	∞	—	B	3	B	B	∞	—	B	1	B
C	0	C	C	1	C	C	∞	—	C	1	C
D	1	D	D	0	D	D	∞	—	D	∞	—
E	∞	—	E	∞	—	E	0	E	E	3	E
F	1	F	F	∞	—	F	3	F	F	0	F

Distance Vector Example 2: Step 2

Optimum 2-hop paths

Table for A			Table for B		
Dst	Cst	Hop	Dst	Cst	Hop
A	0	A	A	4	A
B	4	B	B	0	B
C	7	F	C	2	F
D	7	B	D	3	D
E	2	E	E	4	F
F	5	E	F	1	F

Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	7	F	A	7	B	A	2	A	A	5	B
B	2	F	B	3	B	B	4	F	B	1	B
C	0	C	C	1	C	C	4	F	C	1	C
D	1	D	D	0	D	D	∞	—	D	2	C
E	4	F	E	∞	—	E	0	E	E	3	E
F	1	F	F	2	C	F	3	F	F	0	F




Distance Vector Example 2: Step 3

Optimum 3-hop paths

Table for A			Table for B		
Dst	Cst	Hop	Dst	Cst	Hop
A	0	A	A	4	A
B	4	B	B	0	B
C	6	E	C	2	F
D	7	B	D	3	D
E	2	E	E	4	F
F	5	E	F	1	F

Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	6	F	A	7	B	A	2	A	A	5	B
B	2	F	B	3	B	B	4	F	B	1	B
C	0	C	C	1	C	C	4	F	C	1	C
D	1	D	D	0	D	D	5	F	D	2	C
E	4	F	E	5	C	E	0	E	E	3	E
F	1	F	F	2	C	F	3	F	F	0	F



Comparison of Protocols

Link State

- Knowledge of every router's links (entire graph)
- Every router has $O(\# \text{ edges})$
- Trust a peer's info, do routing computation yourself
- Use Dijkstra's algorithm
- Send updates on any link-state changes
- Ex: OSPF, IS-IS
- Adv: Fast to react to changes

Distance Vector

- Knowledge of neighbors' distance to destinations
- Every router has $O(\# \text{ neighbors} * \# \text{ nodes})$
- Trust a peer's routing computation
- Use Bellman-Ford algorithm
- Send updates periodically or routing decision change
- Ex: RIP, IGRP
- Adv: Less info & lower computational overhead

Similarities of LS and DV Routing

- Shortest-path routing
 - Metric-based, using link weights
 - Routers share a common view of how good a path is
- As such, commonly used *inside* an organization
 - RIP and OSPF are mostly used as *intra*-domain protocols
 - E.g., A small business typically uses RIP, and AT&T uses OSPF
- But the Internet is a “network of networks”
 - How to stitch the many networks together?
 - When networks may not have common goals
 - ... and may not want to share information

Path-Vector Routing

Link-State Routing is Problematic


- Topology information is flooded
 - High bandwidth and storage overhead
 - Forces nodes to divulge sensitive information
- Entire path computed locally per node
 - High processing overhead in a large network
- Minimizes some notion of total distance
 - Works only if policy is shared and uniform
- Typically used only inside an AS
 - E.g., OSPF and IS-IS

Distance Vector is on the Right Track

- **Advantages**
 - Hides details of the network topology
 - Nodes determine only “next hop” toward the dest
- **Disadvantages**
 - Minimizes some notion of total distance, which is difficult in an interdomain setting
 - Slow convergence due to the counting-to-infinity problem (“bad news travels slowly”)
- **Idea: extend the notion of a distance vector**
 - To make it easier to detect loops

Path-Vector Routing


- Extension of distance-vector routing
 - Support flexible routing policies
- Key idea: advertise the entire path
 - Distance vector: send *distance metric* per dest d
 - Path vector: send the *entire path* for each dest d



Used in BGP

Faster Loop Detection


- Node can easily detect a loop
 - Look for its own node identifier in the path
 - E.g., node 1 sees itself in the path “3, 2, 1”
- Node can simply discard paths with loops
 - E.g., node 1 simply discards the advertisement




Path-Vector: Flexible Policies

- Each node can apply local policies
 - Path selection: Which path to use?
 - Path export: Which paths to advertise?

**Node 2 prefers
“2, 3, 1” over “2, 1”**




**Node 1 doesn't let 3
hear the path “1, 2”**



End-to-End Signaling


- Establish end-to-end path in advance
 - Learn the topology (as in link-state routing)
 - End host or router computes and signals a path
 - Signaling: install entry for each circuit at each hop
 - Forwarding: look up the circuit id in the table



Used in MPLS with RSVP

MPLS Overview


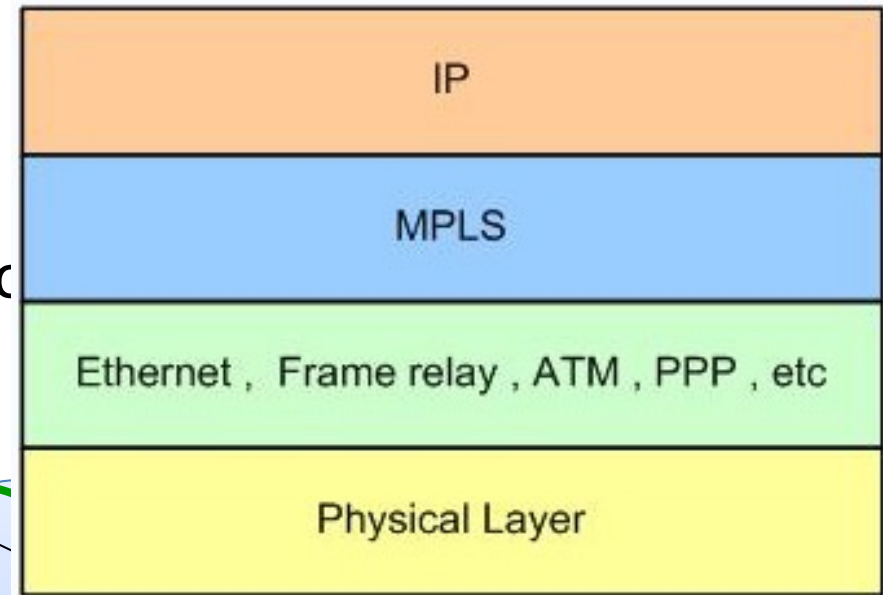
- **Main idea: Virtual circuit**
 - Packets forwarded based only on circuit identifier



Router can forward traffic to the same destination on different interfaces/paths.


MPLS Overview

- **Main idea: Virtual circuit**
 - Packets forwarded based on



Router can forward traffic to the same destination on different interfaces/paths.

Circuit Abstraction: Label Swapping



- **Label-switched paths:** Paths “named” by label at ingress
- **At each hop, MPLS routers:**
 - Use label to determine outgoing interface, new label
 - Thus, push/pop/swap MPLS headers that encapsulate IP
- **Label distribution protocol:** disseminate signaling info
- **Initially from concern with longest-prefix-match speed**
 - Now use in other applications, e.g., intra-AS traffic management

Conclusions

- Distance-vector routing
 - Pro: Less information and computation than link state
 - Con: Slower convergence (e.g., count to infinity)
- Path-vector routing
 - Share entire path, not distance: faster convergence
 - More flexibility in selecting paths
- Different goals / metrics if inter- or intra-domain