


Routing Convergence

Note: The slides are adapted from the materials from Prof. Richard Han at CU Boulder and Profs. Jennifer Rexford and Mike Freedman at Princeton University, and the networking book (Computer Networking: A Top Down Approach) from Kurose and Ross.

Routing Changes



- **Topology changes:** new route to the same place
- **Host mobility:** route to a different place

Topology Changes


Two Types of Topology Changes

- **Planned**

- Maintenance: shut down a node or link
- Energy savings: shut down a node or link
- Traffic engineering: change routing configuration

- **Unplanned Failures**


- Fiber cut,
faulty equipment,
power outage,
software bugs, ...



Detecting Topology Changes

- **Beaconing**

- Periodic “hello” messages in both directions
- Detect a failure after a few missed “hellos”




- **Performance trade-offs**

- Detection delay
- Overhead on link bandwidth and CPU
- Likelihood of false detection

Routing Convergence: Link-State Routing


Convergence

- **Control plane**
 - All nodes have consistent information
- **Data plane**
 - All nodes forward packets in a consistent way





Transient Disruptions

- **Detection delay**
 - A node does not detect a failed link immediately
 - ... and forwards data packets into a “blackhole”
 - Depends on timeout for detecting lost hellos



Transient Disruptions

- **Inconsistent link-state database**
 - Some routers know about failure before others
 - Inconsistent paths cause transient forwarding loops



Convergence Delay

- Sources of convergence delay
 - Detection latency
 - Updating control-plane information
 - Computing and install new forwarding tables
- Performance during convergence period
 - Lost packets due to blackholes and TTL expiry
 - Looping packets consuming resources
 - Out-of-order packets reaching the destination
- Very bad for VoIP, online gaming, and video

Reducing Convergence Delay


- **Faster detection**
 - Smaller hello timers, better link-layer technologies
- **Faster control plane**
 - Flooding immediately
 - Sending routing messages with high-priority
- **Faster computation**
 - Faster processors, and incremental computation
- **Faster forwarding-table update**
 - Data structures supporting incremental updates

Slow Convergence in Distance-Vector Routing

Distance Vector: Link Cost Changes

- Link cost decreases and recovery

- Node updates the distance table
- If cost change in least cost path, notify neighbors



D^Y = Distances known to Y

D^Y		via	
	X	Z	
X	(4)	6	

D^Z		via	
	X	Y	
X	50	(5)	




Distance Vector: Link Cost Changes

- Link cost decreases and recovery

- Node updates the distance table

- If cost change in least cost path, notify neighbors



D^Y = Distances known to Y

D^Y		
	X	Z
via		
X	4	6

D^Y		
	X	Z
	1	6

D^Z		
	X	Y
via		
X	50	5

D^Z		
	X	Y
	50	5

$c(X,Y)$
change





Distance Vector: Link Cost Changes

- Link cost decreases and recovery

- Node updates the distance table


- If cost change in least cost path, notify neighbors



Distance Vector: Link Cost Changes

- Link cost increases and failures

- Bad news travels slowly
- “Count to infinity” problem!



Y		
D	X	Z
X	4	6

X		
D	X	Z
X	60	6

Z		
D	X	Y
X	50	5



Z		
D	X	Y
X	50	5



Distance Vector: Link Cost Changes

- Link cost increases and failures

- Bad news travels slowly
- “Count to infinity” problem!





Distance Vector: Poison Reverse

- If Z routes through Y to get to X :

- Z tells Y its (Z's) distance to X is infinite (so Y won't route to X via Z)

- Still, can have problems in larger networks




Redefining Infinity

- Avoid “counting to infinity”
 - By making “infinity” smaller!
- Routing Information Protocol (RIP)
 - All links have cost 1
 - Valid path distances of 1 through 15
 - ... with 16 representing infinity
- Used mainly in small networks

Reducing Convergence Time With Path-Vector Routing (e.g., Border Gateway Protocol)


Path-Vector Routing

- Extension of distance-vector routing
 - Support flexible routing policies
 - Avoid count-to-infinity problem
- Key idea: advertise the entire path
 - Distance vector: send distance metric per dest d
 - Path vector: send the entire path for each dest d



Faster Loop Detection

- Node can easily detect a loop
 - Look for its own node identifier in the path
 - E.g., node 1 sees itself in the path “3, 2, 1”
- Node can simply discard paths with loops
 - E.g., node 1 simply discards the advertisement




Causes of BGP Routing Changes



- **Topology changes**
 - Equipment going up or down
 - Deployment of new routers or sessions
- **BGP session failures**
 - Due to equipment failures, maintenance, etc.
 - Or, due to congestion on the physical path
- **Changes in routing policy**
 - Changes in preferences in the routes
 - Changes in whether the route is exported
- **Persistent protocol oscillation**
 - Conflicts between policies in different ASes

BGP Session Failure

- **BGP runs over TCP**
 - BGP only sends updates when changes occur
 - TCP doesn't detect lost connectivity on its own
- **Detecting a failure**
 - Keep-alive: 60 seconds
 - Hold timer: 180 seconds
- **Reacting to a failure**
 - Discard all routes learned from neighbor
 - Send new updates for any routes that change



Routing Change: Before and After




Routing Change: Path Exploration

- **AS 1**

- Delete the route (1,0)
- Switch to next route (1,2,0)
- Send route (1,2,0) to AS 3

- **AS 3**


- Sees (1,2,0) replace (1,0)
- Compares to route (2,0)
- Switches to using AS 2



Routing Change: Path Exploration

- Initial: All AS use direct
- Then destination 0 dies
 - All ASes lose direct path
 - All switch to longer paths
 - Eventually withdrawn
- How many intermediate routes following (2,0) withdrawal until no route known to 2?

$(2,0) \rightarrow (2,1,0) \rightarrow (2,3,0) \rightarrow (2,1,3,0) \rightarrow \text{null}$




BGP Converges Slowly

- Path vector avoids count-to-infinity
 - But, ASes still must explore many alternate paths to find highest-ranked available path
- Fortunately, in practice
 - Most popular destinations have stable BGP routes
 - Most instability lies in a few unpopular destinations
- Still, lower BGP convergence delay is a goal
 - Can be tens of seconds to tens of minutes

BGP Instability

Stable Paths Problem (SPP) Instance

- **Node**
 - BGP-speaking router
 - Node 0 is destination
- **Edge**
 - BGP adjacency
- **Permitted paths**
 - Set of routes to 0 at each node
 - Ranking of the paths




Stable Paths Problem (SPP) Instance

- 1 will use a direct path to 0


(A) True (B) False

- 5 has a path to 0

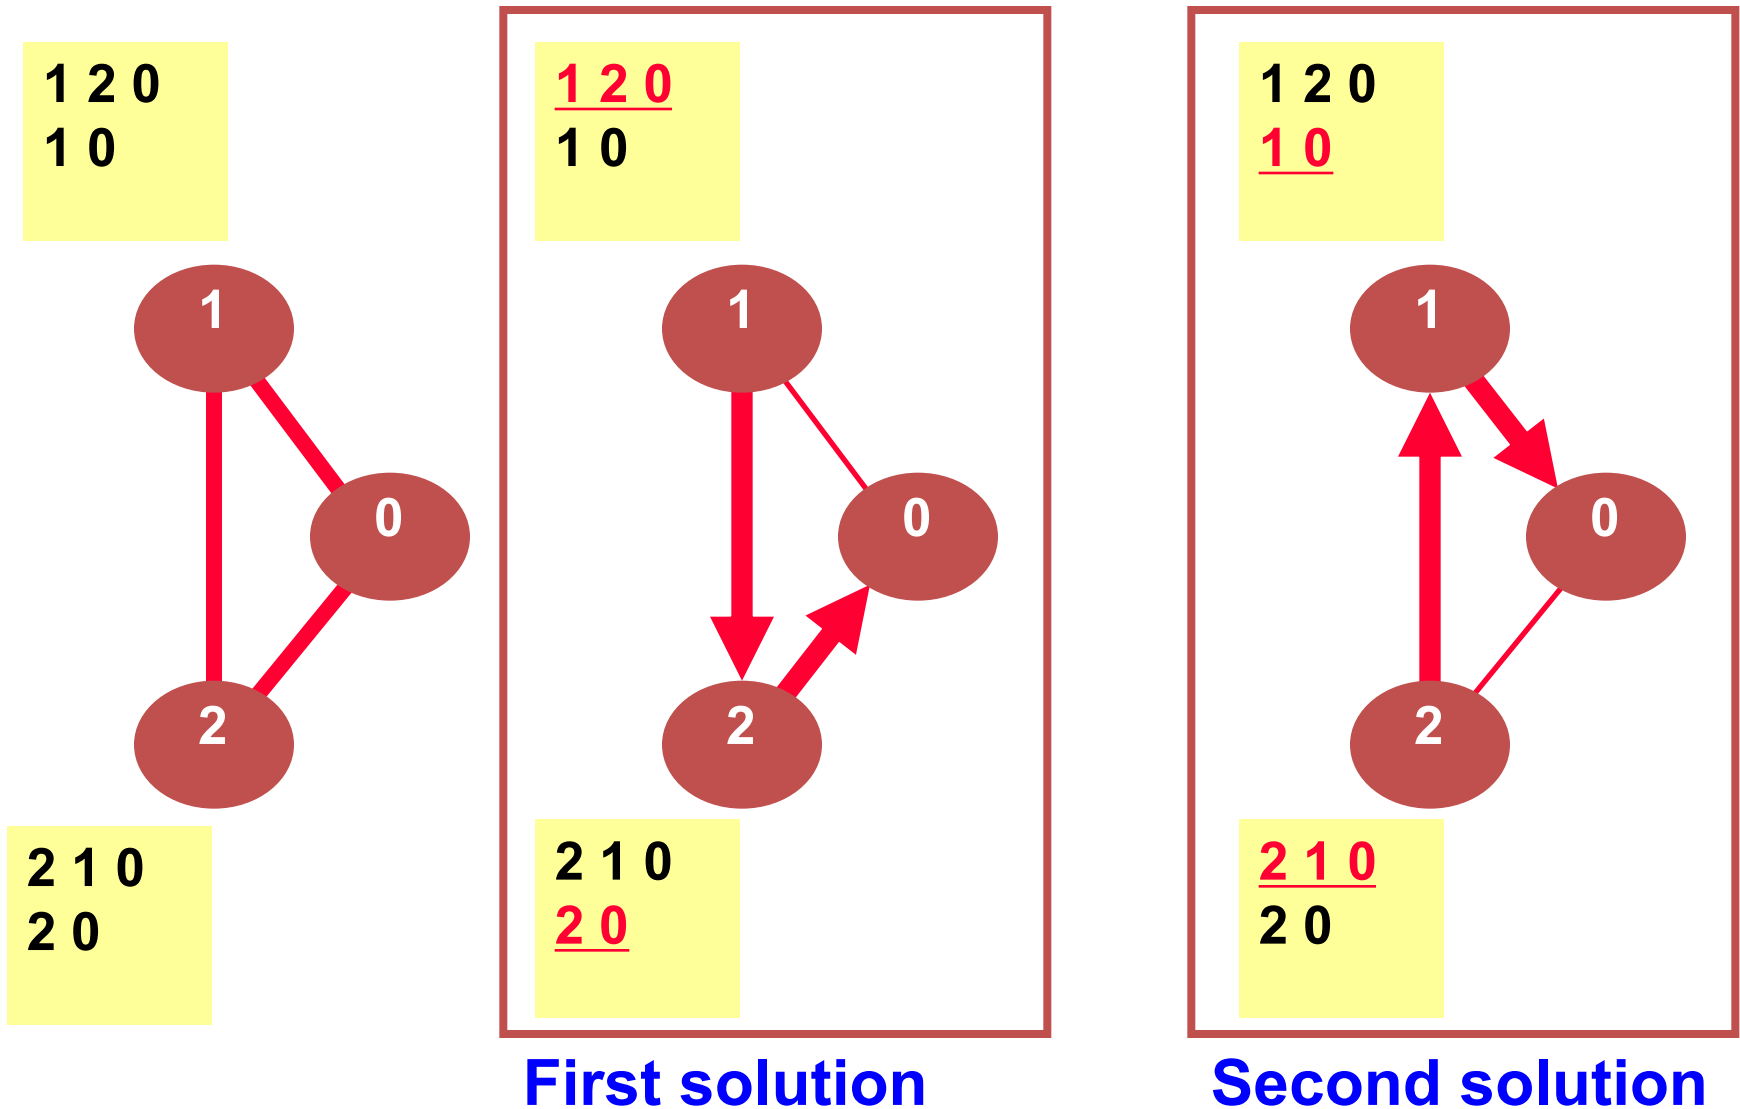
(A) True (B) False




Stable Paths Problem (SPP) Instance



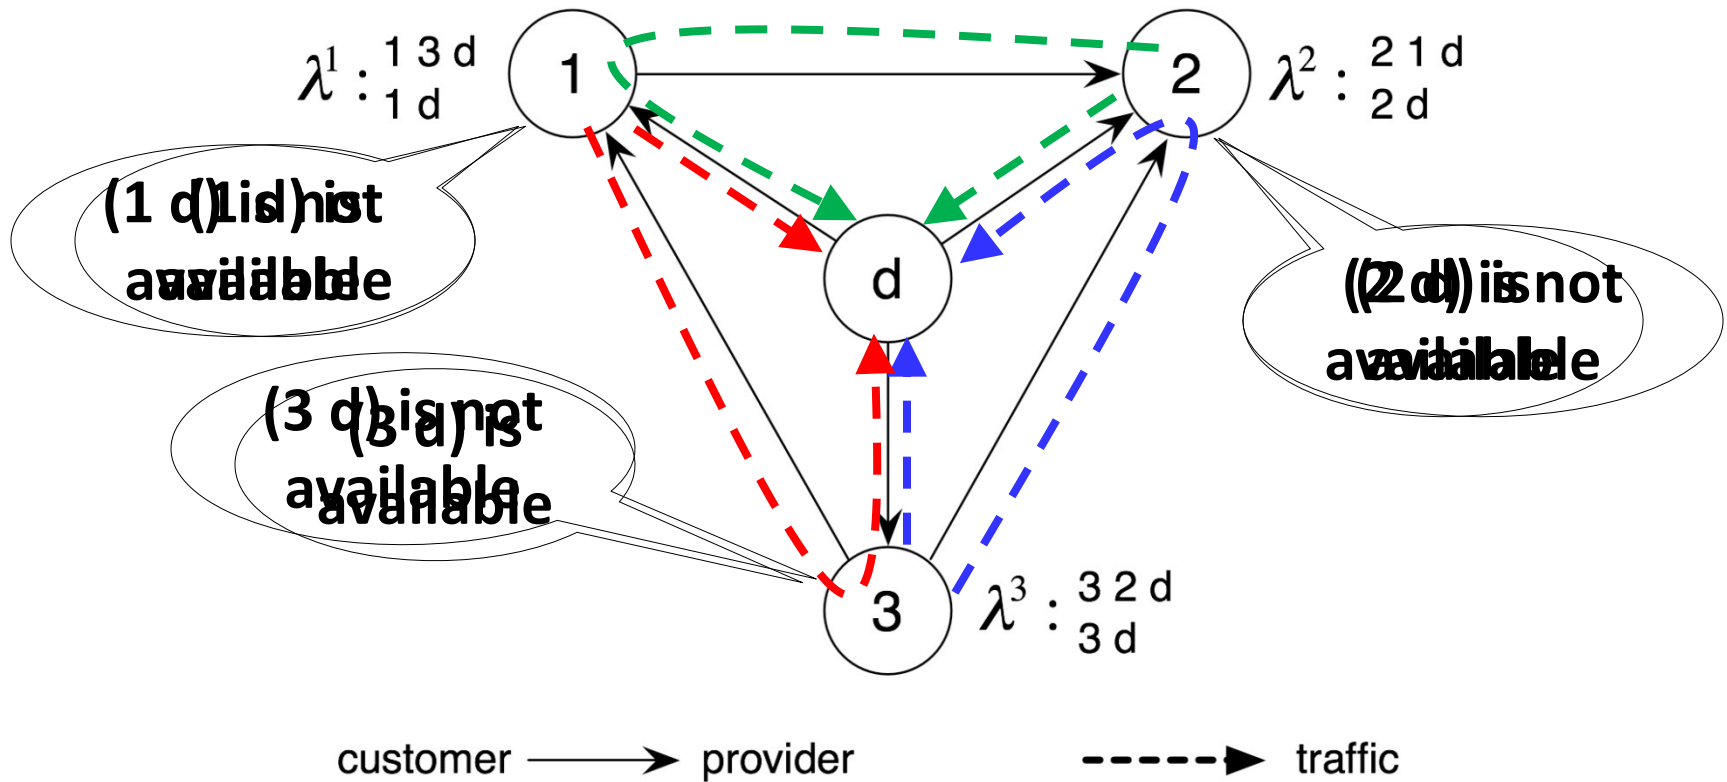
SPP May Have Multiple Solutions



An SPP May Have No Solution



BGP Not Guaranteed to Converge



Example known as a “dispute wheel”

Avoiding BGP Instability

- Detecting conflicting policies
 - Computationally expensive
 - Requires too much cooperation
- Detecting oscillations
 - Observing the repetitive BGP routing messages
- Restricted routing policies and topologies
 - Policies based on business relationships

Conclusion

- The only constant is change
 - Planned topology and configuration changes
 - Unplanned failure and recovery
- Routing-protocol convergence
 - Transient period of disagreement
 - Blackholes, loops, and out-of-order packets
- Routing instability
 - Permanent conflicts in routing policy
 - Leading to bi-stability or oscillation