

Hashing in Networked Systems

Note: The slides are adapted from the materials from Prof. Richard Han at CU Boulder and Profs. Jennifer Rexford and Mike Freedman at Princeton University, and the networking book (Computer Networking: A Top Down Approach) from Kurose and Ross.


Hashing

- Hash function

- Function that maps a large, possibly variable-sized datum into a small datum, often a single integer that serves to index an associative array
- In short: maps n -bit datum into k buckets ($k \ll 2^n$)
- Provides time- & space-saving data structure for lookup

- Main goals:

- Low cost
- Deterministic
- Uniformity (load balanced)




Today's outline

- Uses of hashing
 - Equal-cost multipath routing in switches
 - Network load balancing in server clusters
 - Per-flow statistics in switches (QoS, IDS)
 - Caching in cooperative CDNs and P2P file sharing
 - Data partitioning in distributed storage services
- Various hashing strategies
 - Modulo hashing
 - Consistent hashing
 - Bloom Filters


Uses of Hashing

Equal-cost multipath routing (ECMP)




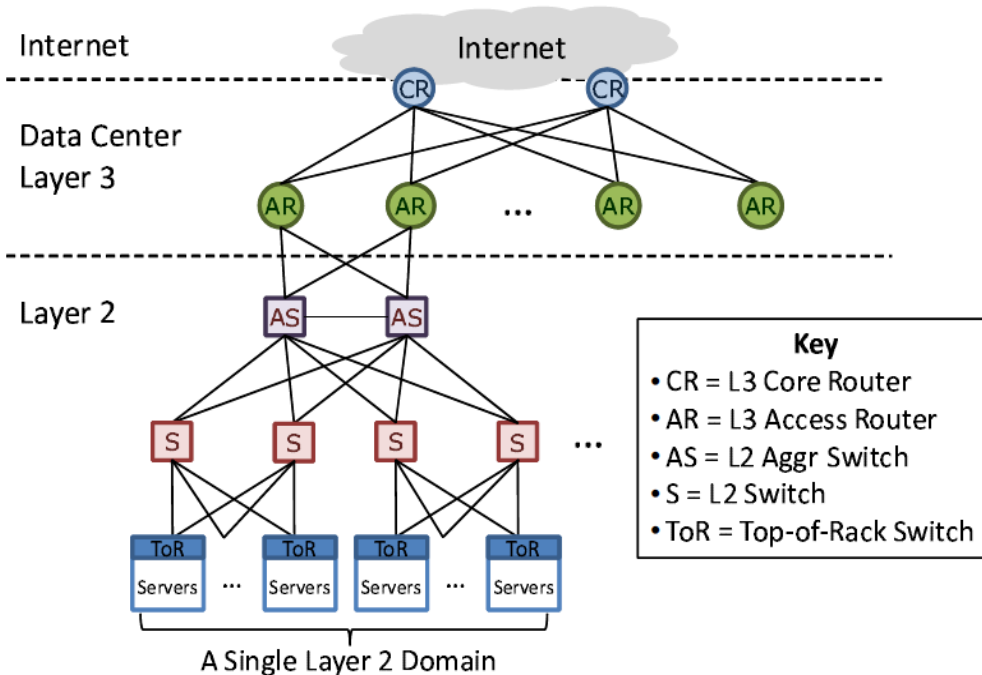
- **ECMP**
 - Multipath routing strategy that splits traffic over multiple paths for load balancing
- **Why not just round-robin packets?**
 - Reordering (lead to triple duplicate ACK in TCP?)
 - Different RTT per path (for TCP RTO)...
 - Different MTUs per path

Equal-cost multipath routing (ECMP)




- Path-selection via hashing
 - # buckets = # outgoing links
 - Hash network information (source/dest IP addrs) to select outgoing link: preserves flow affinity

Now: ECMP in datacenters




- **Datacenter networks are multi-rooted tree**
 - Goal: Support for 100,000s of servers
 - Recall Ethernet spanning tree problems: No loops
 - L3 routing and ECMP: Take advantage of multiple paths

Network load balancing



- **Goal: Split requests evenly over k servers**
 - Map new flows to any server
 - Packets of existing flows continue to use same server
- **3 approaches**
 - Load balancer terminates TCP, opens own connection to server
 - Virtual IP / Dedicated IP (VIP/DIP) approaches
 - One global-facing virtual IP represents all servers in cluster
 - Hash client's network information (source IP:port)
 - **NAT approach:** Replace virtual IP with server's actual IP
 - **Direct Server Return (DSR)**

Load balancing with DSR




- Servers bind to both virtual and dedicated IP
- Load balancer just replaces dest MAC addr
- Server sees client IP, responds directly
 - Packet in reverse direction do not pass through load balancer
 - Greater scalability, particularly for traffic with asymmetric bandwidth (e.g., HTTP GETs)

Per-flow state in switches


- Switches often need to maintain connection records or per-flow state
 - Quality-of-service for flows
 - Flow-based measurement and monitoring
 - Payload analysis in Intrusion Detection Systems (IDSs)
- On packet receipt:
 - Hash flow information (packet 5-tuple)
 - Perform lookup if packet belongs to known flow
 - Otherwise, possibly create new flow entry
 - Probabilistic match (false positives) may be okay

Cooperative Web CDNs

- Tree-like topology of cooperative web caches
 - Check local
 - If miss, check siblings / parent
- One approach
 - Internet Cache Protocol (ICP)
 - UDP-based lookup, short timeout
- Alternative approach
 - A priori guess is siblings/children have content
 - Nodes share hash table of cached content with parent / siblings
 - Probabilistic check (false positives) okay, as actual ICP lookup to neighbor could just return false




Hash tables in P2P file-sharing



- Two-layer network (e.g., Gnutella, Kazaa)
 - Ultrapeers are more stable, not NATted, higher bandwidth
 - Leaf nodes connect with 1 or more ultrapeers
- Ultrapeers handle content searchers
 - Leaf nodes send hash table of content to ultrapeers
 - Search requests flooded through ultrapeer network
 - When ultrapeer gets request, checks hash tables of its children for match

Data partitioning

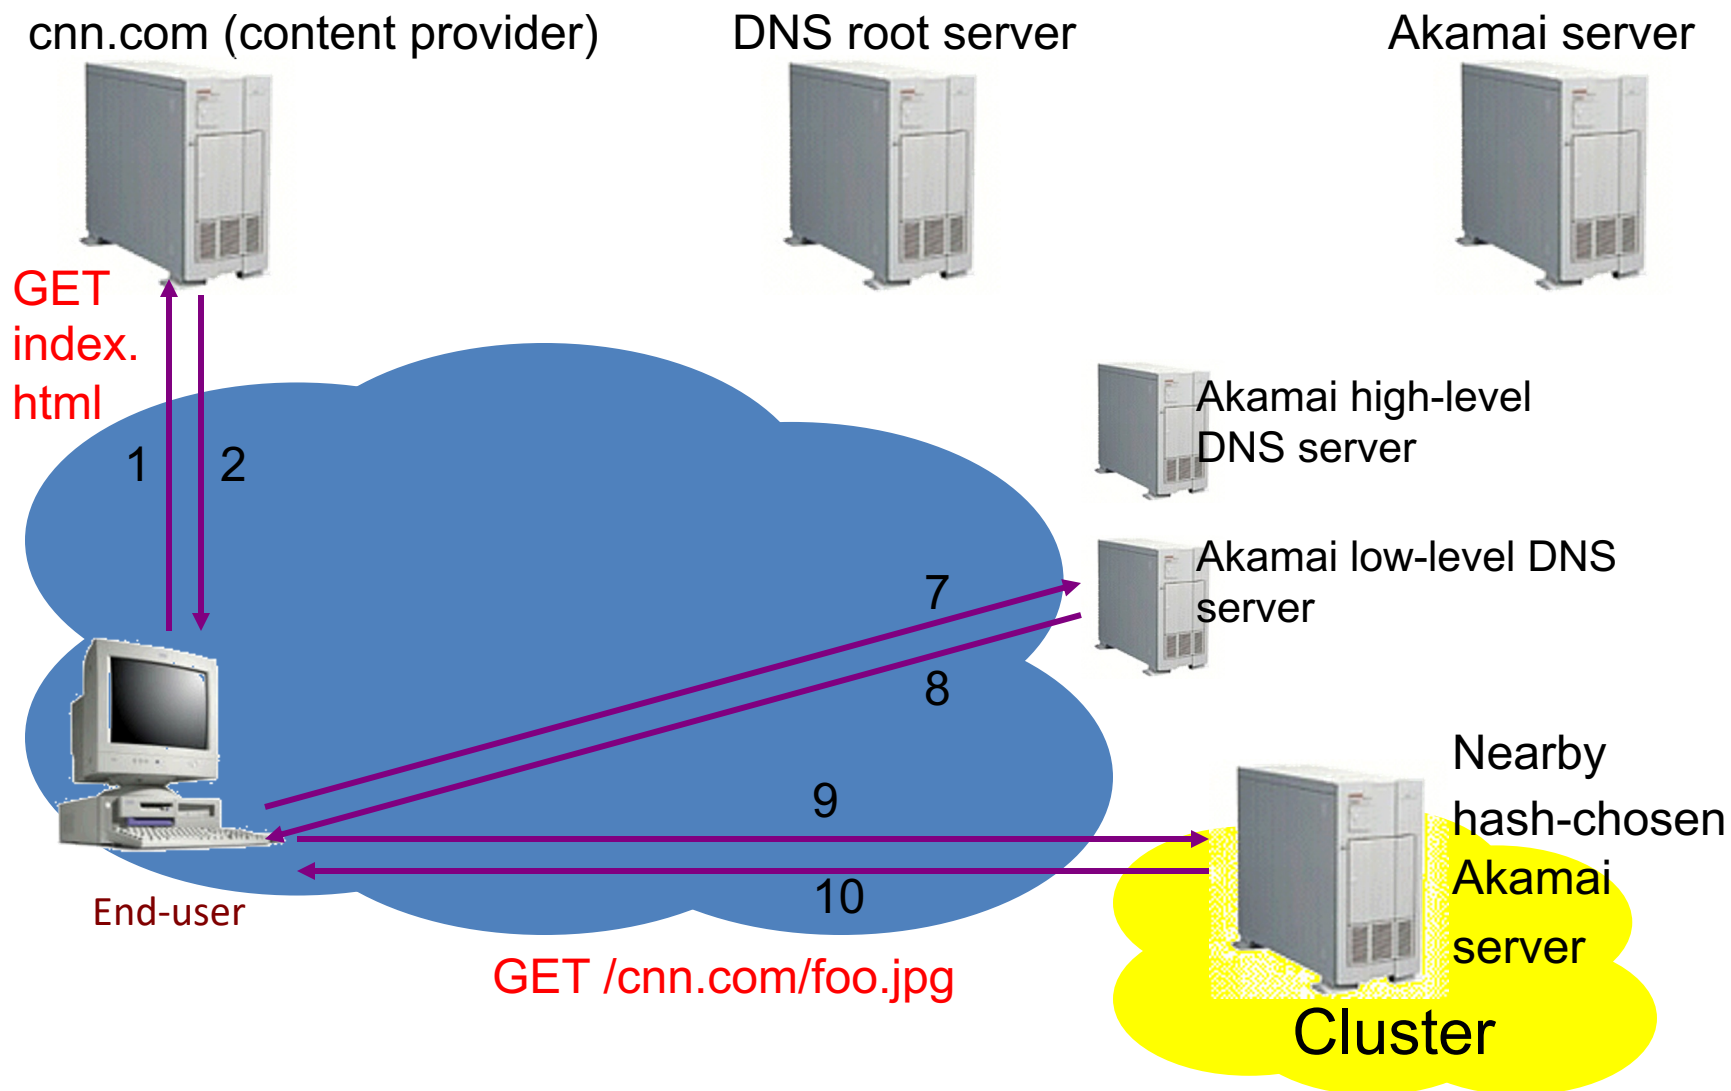


- Network load balancing: All machines are equal
- Data partitioning: Machines store different content
- Non-hash-based solution
 - “Directory” server maintains mapping from $O(\text{entries})$ to machines (e.g., Network file system, Google File System)
 - Named data can be placed on any machine
- Hash-based solution
 - Nodes maintain mappings from $O(\text{buckets})$ to machines
 - Data placed on the machine that owns the name's bucket

Examples of data partitioning

- **Akamai**
 - 1000 clusters around Internet, each ≥ 1 servers
 - Hash (URL's domain) to map to one server
 - Akamai DNS aware of hash function, returns machine that
 1. is in geographically-nearby cluster
 2. manages particular customer domain
- **Memcached (Facebook, Twitter, ...)**
 - Employ k machines for in-memory key-value caching
 - On read:
 - Check memcache
 - If miss, read data from DB, write to memcache
 - On write: invalidate cache, write data to DB

How Akamai Works – Already Cached



Hashing Techniques


Basic Hash Techniques

- Simple approach for uniform data
 - If data distributed uniformly over N , for $N \gg n$
 - Hash fn = $\langle \text{data} \rangle \bmod n$
 - Fails goal of uniformity if data not uniform
- Non-uniform data, variable-length strings
 - Typically split strings into blocks
 - Perform rolling computation over blocks
 - CRC32 checksum
 - Cryptographic hash functions (SHA-1 has 64 byte blocks)

Applying Basic Hashing


- Consider problem of data partition:
 - Given document X , choose one of k servers to use
- Suppose we use modulo hashing
 - Number servers $1..k$
 - Place X on server $i = (X \bmod k)$
 - Problem? Data may not be uniformly distributed
 - Place X on server $i = \text{hash}(X) \bmod k$
 - Problem?
 - What happens if a server fails or joins ($k \rightarrow k \pm 1$)?
 - What is different clients has different estimate of k ?
 - Answer: All entries get remapped to new nodes!

Consistent Hashing



- Consistent hashing partitions key-space among nodes
- Contact appropriate node to lookup/store key
 - Blue node determines red node is responsible for key₁
 - Blue node sends lookup or insert to red node

Consistent Hashing




- Partitioning key-space among nodes
 - Nodes choose random identifiers: e.g., `hash(IP)`
 - Keys randomly distributed in ID-space: e.g., `hash(URL)`
 - Keys assigned to node “nearest” in ID-space
 - Spreads ownership of keys evenly across nodes

Consistent Hashing

- **Construction**

- Assign n hash buckets to random points on mod 2^k circle; hash key size = k
- Map object to random position on circle
- Hash of object = closest clockwise bucket
 - *successor* (key) \rightarrow bucket




- **Desired features**

- **Balanced**: No bucket has disproportionate number of objects
- **Smoothness**: Addition/removal of bucket does not cause movement among existing buckets (only immediate buckets)
- **Spread and load**: Small set of buckets that lie near object


Consistent hashing and failures

- Consider network of n nodes
- If each node has 1 bucket
 - Owns $1/n^{\text{th}}$ of keypace *in expectation*
- If a node fails:
 - (A) Nobody owns keypace (B) Keyspace assigned to random node
 - (C) Successor owns keyspaces (D) Predecessor owns keyspace
- After a node fails:
 - (A) Load is equally balanced over all nodes
 - (B) Some node has disproportional load compared to others







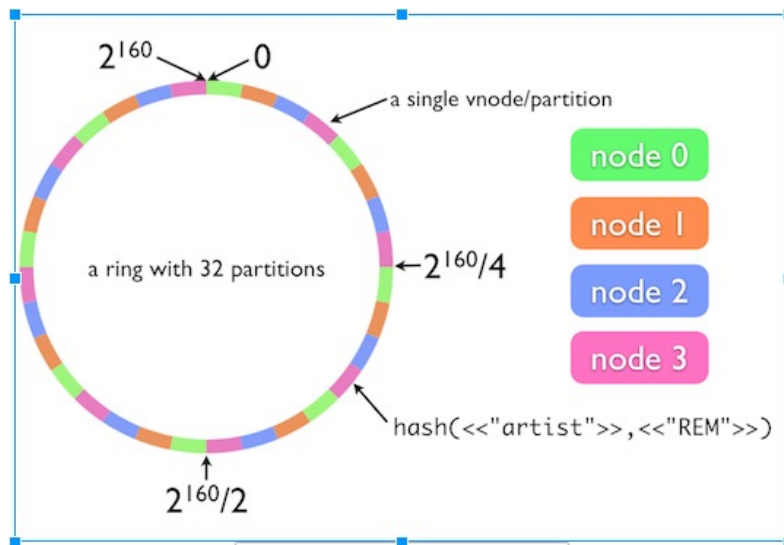
Consistent hashing and failures

- Consider network of n nodes
- If each node has 1 bucket
 - Owns $1/n^{\text{th}}$ of keyspace *in expectation*
- If a node fails:
 - Its *successor* takes over bucket
 - Achieves smoothness goal: Only localized shift, not $O(n)$
 - But now successor owns 2 buckets: keyspace of size $2/n$
- Instead, if each node maintains v random nodeIDs, not 1
 - “Virtual” nodes spread over ID space, each of size $1/vn$
 - Upon failure, v successors take over, each now stores $(v+1)/vn$



Example: Cassandra


- Cassandra adopts consistent hashing with virtual nodes for data partitioning



Consistent hashing vs. DHTs


	Consistent Hashing	Distributed Hash Tables
Routing table size	$O(n)$	$O(\log n)$
Lookup / Routing	$O(1)$	$O(\log n)$
Join/leave: Routing updates	$O(n)$	$O(\log n)$
Join/leave: Key Movement	$O(1)$	$O(1)$

Distributed Hash Table




- Nodes' neighbors selected from particular distribution
 - Visual keyspace as a tree in distance from a node

Distributed Hash Table



- Nodes' neighbors selected from particular distribution
 - Visual keyspace as a tree in distance from a node
 - At least one neighbor known per subtree of increasing size /distance from node


Distributed Hash Table



- Nodes' neighbors selected from particular distribution
 - Visual keypace as a tree in distance from a node
 - At least one neighbor known per subtree of increasing size /distance from node
- Route greedily towards desired key via overlay hops


The Chord DHT

- Chord ring: ID space mod 2^{160}
 - $nodeid = SHA1(IP\ address, i)$
for $i=1..v$ virtual IDs
 - $keyid = SHA1(name)$
- Routing correctness:
 - Each node knows successor and predecessor on ring
- Routing efficiency:
 - Each node knows $O(\log n)$ well-distributed neighbors



Basic lookup in Chord

```
lookup (id):  
    if ( id > pred.id &&  
        id <= my.id )  
        return my.id;  
    else  
        return succ.lookup(id);
```




- Route hop by hop via successors
 - $O(n)$ hops to find destination id

Efficient lookup in Chord

```

lookup (id):
    if ( id > pred.id &&
        id <= my.id )
return my.id;
else
    // fingers() by decreasing distance
    for finger in fingers():
        if id >= finger.id
            return finger.lookup(id);
    return succ.lookup(id);


```



- Route greedily via distant “finger” nodes
 - $O(\log n)$ hops to find destination id


Building routing tables

Routing Tables




For i in $0 \dots \log n$:

$$\text{finger}[i] = \text{successor} ((\text{my.id} + 2^i) \bmod 2^{160})$$




Joining and managing routing

- **Join:**
 - Choose nodeid
 - *Lookup* (*my.id*) to find place on ring
 - During lookup, discover future successor
 - Learn predecessor from successor
 - Update succ and pred that you joined
 - Find fingers by *lookup* ($(\text{my.id} + 2^i) \bmod 2^{160}$)
- **Monitor:**
 - If doesn't respond for some time, find new
- **Leave: Just go, already!**
 - (Warn your neighbors if you feel like it)



Performance optimizations



- Routing entries need not be drawn from strict distribution as finger algorithm shown
 - Choose node with lowest latency to you
 - Will still get you $\sim \frac{1}{2}$ closer to destination
- Less flexibility in choice as closer to destination


DHT Design Goals

- An “overlay” network with:
 - Flexible mapping of keys to physical nodes
 - Small network diameter
 - Small degree (fanout)
 - Local routing decisions
 - Robustness to churn
 - Routing flexibility
 - Decent locality (low “stretch”)
- Different “storage” mechanisms considered:
 - Persistence w/ additional mechanisms for fault recovery
 - Best effort caching and maintenance via soft state

Chord DHT Example


- Assume an identifier space $[0..8]$
- Node n1 joins

i	$id+2^i$	succ
0	2	1
1	3	1
2	5	1



Chord DHT Example

- Node n2 joins



i	id+2 ⁱ	succ
0	2	2
1	3	1
2	5	1

i	id+2 ⁱ	succ
0	3	1
1	4	1
2	6	1

Chord Example


- Nodes n0 and n6 join

i	id+2 ⁱ	succ
0	1	1
1	2	2
2	4	6

i	id+2 ⁱ	succ
0	2	2
1	3	6
2	5	6

i	id+2 ⁱ	succ
0	7	0
1	0	0
2	2	2

i	id+2 ⁱ	succ
0	3	6
1	4	6
2	6	6



Chord

- Nodes: n1, n2, n0, n6
- Items: f7, f1

i	id+2 ⁱ	succ
0	1	1
1	2	2
2	4	6

Items

7


i	id+2 ⁱ	succ
0	2	2
1	3	6
2	5	6

Items

1

i	id+2 ⁱ	succ
0	7	0
1	0	0
2	2	2


i	id+2 ⁱ	succ
0	3	6
1	4	6
2	6	6



Chord

i	id+2 ⁱ	succ
0	1	1
1	2	2
2	4	6

Items

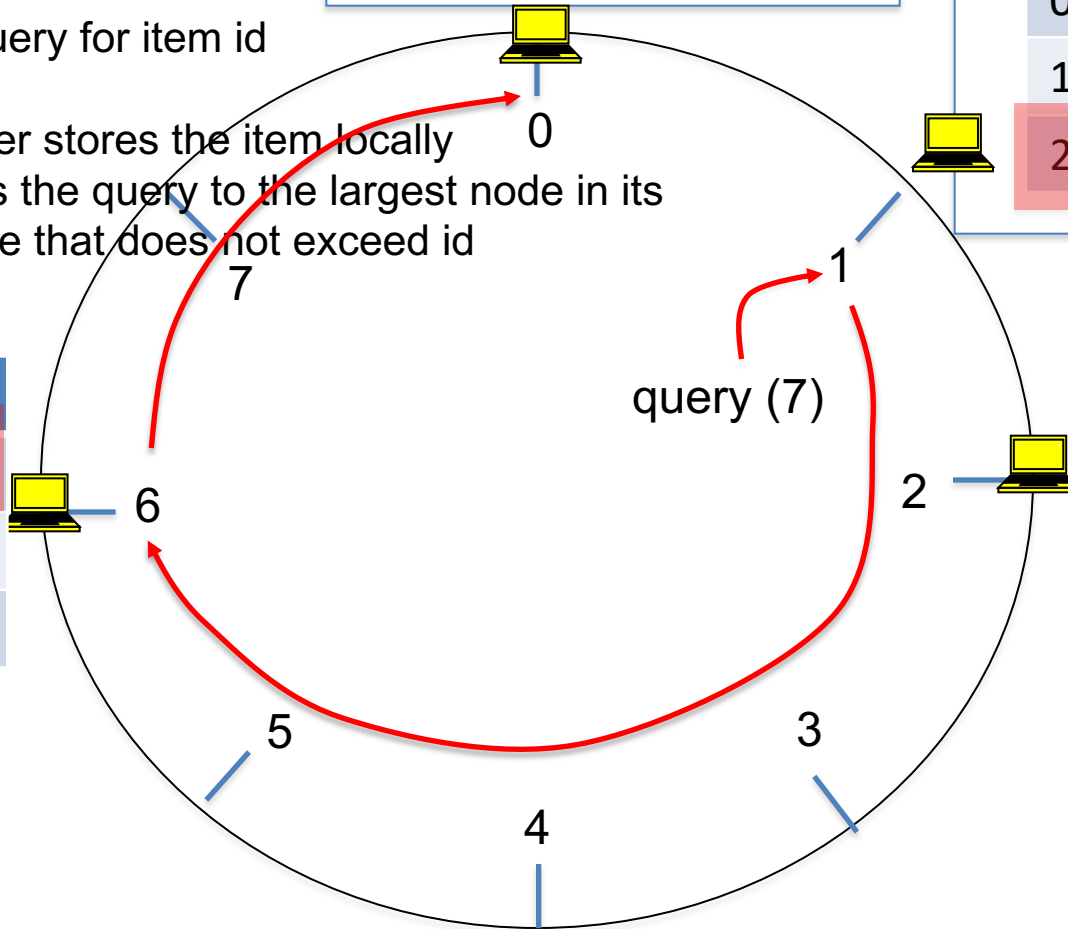


i	id+2 ⁱ	succ	Items
0	2	2	
1	3	6	
2	5	6	1

- Upon receiving a query for item id
- A node:
 - Checks whether stores the item locally
 - If not, forwards the query to the largest node in its successor table that does not exceed id


i	id+2 ⁱ	succ
0	7	0
1	0	0
2	2	2

i	id+2 ⁱ	succ
0	3	6
1	4	6
2	6	6



Bloom Filters

- Data structure for probabilistic membership testing
 - Small amount of space, constant time operations
 - False positives possible, no false negatives
 - Useful in per-flow network statistics, sharing information between cooperative caches, etc.
- Basic idea using hash fn' s and bit array
 - Use k independent hash functions to map item to array
 - If all array elements are 1, it' s present. Otherwise, not



Bloom Filters

Start with an m bit array, filled with 0s.

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

To insert, hash each item k times. If $H_i(x) = a$, set $Array[a] = 1$.

0	1	0	0	1	0	1	0	0	1	1	1	0	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

To check if y is in set, check array at $H_i(y)$. All k values must be 1.

0	1	0	0	1	0	1	0	0	1	1	1	0	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Possible to have a false positive: all k values are 1, but y is not in set.

0	1	0	0	1	0	1	0	0	1	1	1	0	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Summary

- **Peer-to-peer systems**
 - Unstructured systems
 - Finding hay, performing keyword search
 - Structured systems (DHTs)
 - Finding needles, exact match
- **Distributed hash tables**
 - Based around consistent hashing with views of $O(\log n)$
 - Chord, Pastry, CAN, Koorde, Kademlia, Tapestry, Viceroy, ...
- **Lots of systems issues**
 - Heterogeneity, storage models, locality, churn management, underlay issues, ...
 - DHTs deployed in wild: Vuze (Kademlia) has 1M+ active users