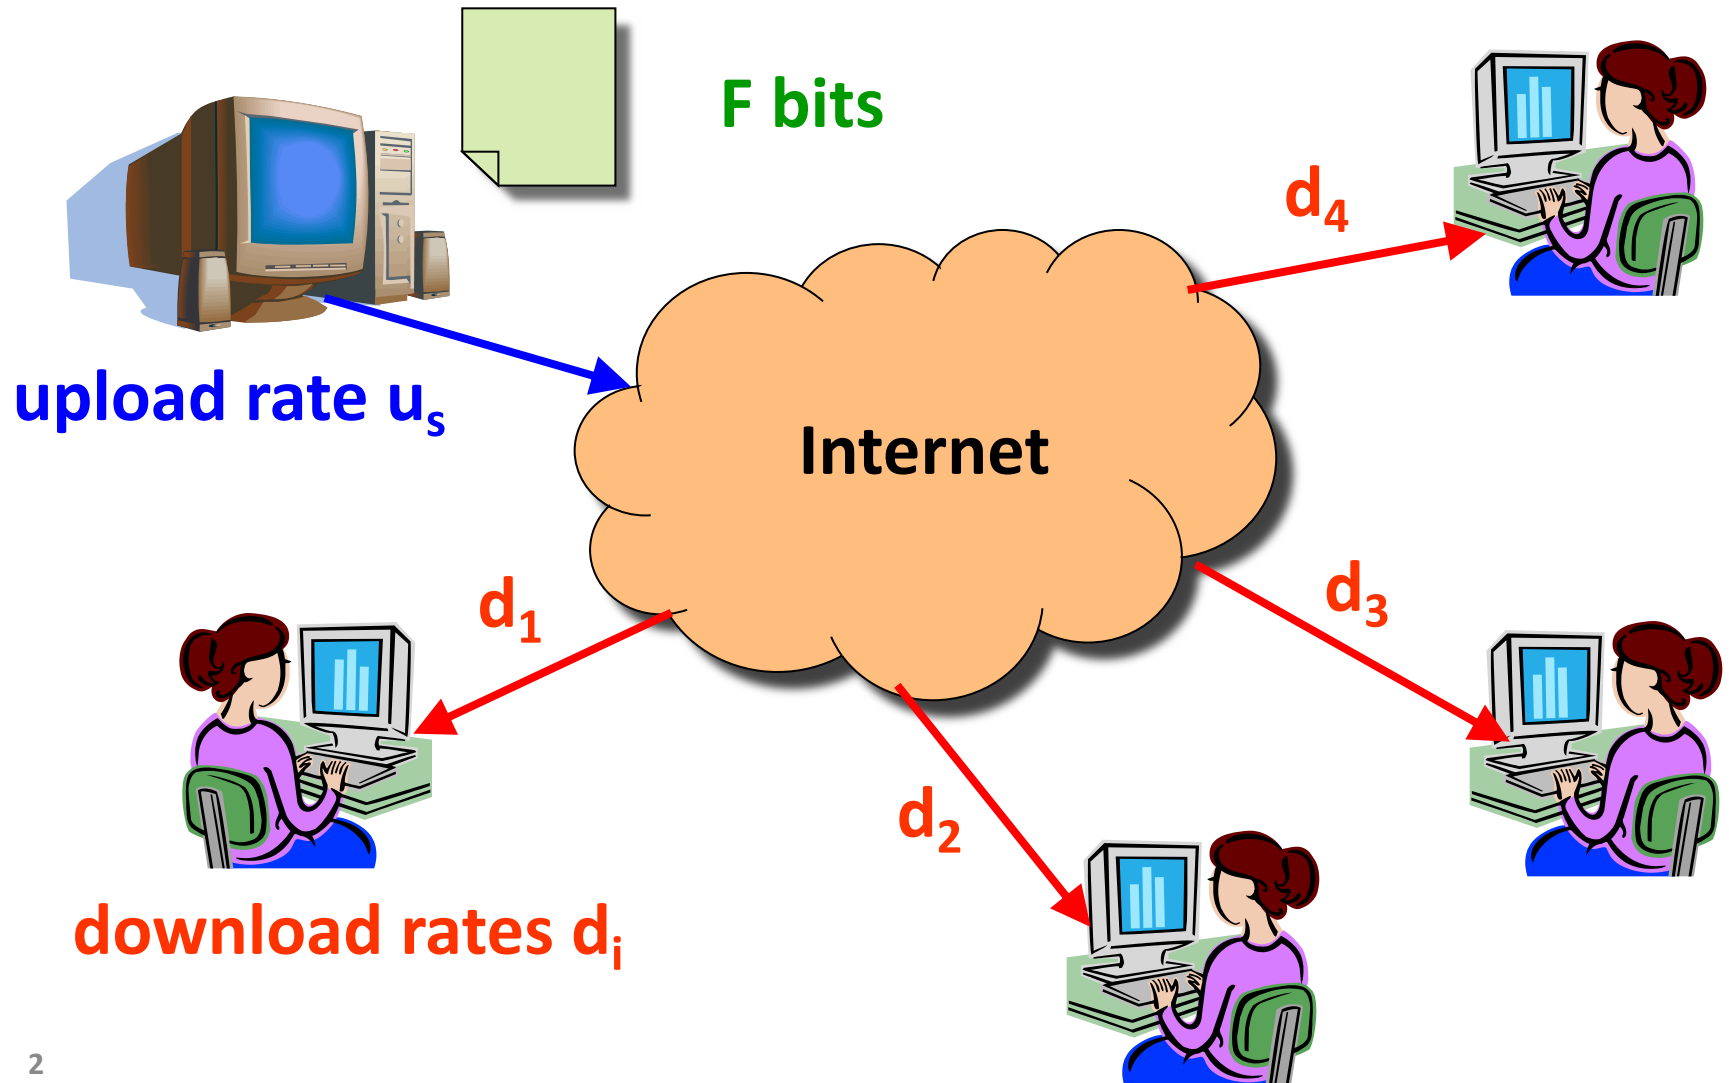




Peer-to-Peer File Sharing

Note: The slides are adapted from the materials from Prof. Richard Han at CU Boulder and Profs. Jennifer Rexford and Mike Freedman at Princeton University, and the networking book (Computer Networking: A Top Down Approach) from Kurose and Ross.

Server Distributing a Large File



Server Distributing a Large File

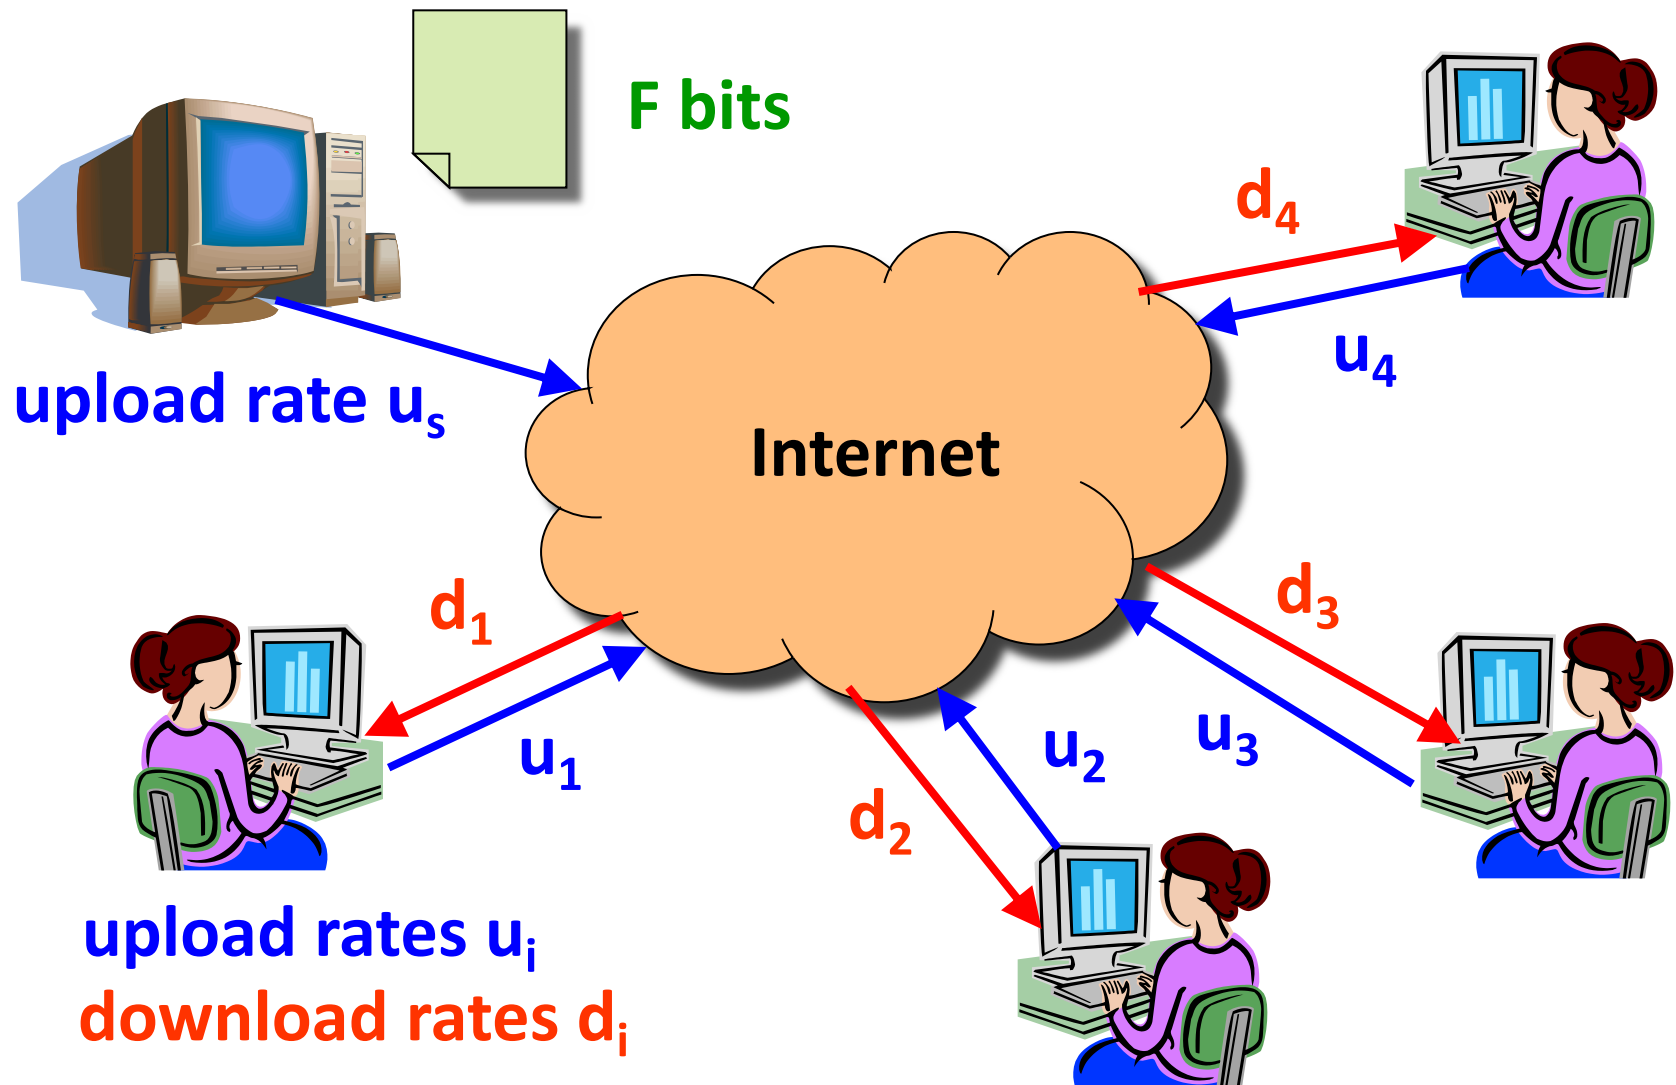
- Sending an F -bit file to N receivers
 - Transmitting NF bits at rate u_s
 - ... takes at least NF/u_s time
- Receiving the data at the slowest receiver
 - Slowest receiver has download rate $d_{min} = \min_i \{d_i\}$
 - ... takes at least F/d_{min} time
- Download time: $\max\{NF/u_s, F/d_{min}\}$

Speeding Up the File Distribution

- Increase the server upload rate
 - Higher link bandwidth at the server
 - Multiple servers, each with their own link
- Alternative: have the receivers help
 - Receivers get a copy of the data
 - ... and redistribute to other receivers
 - To reduce the burden on the server



Peers Help Distributing a Large File



Peers Help Distributing a Large File

- Components of distribution latency
 - Server must send each bit: min time F/u_s
 - Slowest peer must receive each bit: min time F/d_{min}
- Upload time using all upload resources
 - Total number of bits: NF
 - Total upload bandwidth $u_s + \sum_i(u_i)$
- Total: $\max\{F/u_s, F/d_{min}, NF/(u_s + \sum_i(u_i))\}$

Peer-to-Peer is Self-Scaling

- Download time grows slowly with N
 - Client-server: $\max\{NF/u_s, F/d_{min}\}$
 - Peer-to-peer: $\max\{F/u_s, F/d_{min}, NF/(u_s + \sum_i(u_i))\}$
- But...
 - Peers may come and go
 - Peers need to find each other
 - Peers need to be willing to help each other

Locating the Relevant Peers

- Three main approaches
 - Central directory (Napster)
 - Query flooding (Gnutella)
 - Hierarchical overlay (Kazaa, modern Gnutella)
- Design goals
 - Scalability
 - Simplicity
 - Robustness
 - Plausible deniability

Peer-to-Peer Networks: Napster

- **Napster history: the rise**
 - 1/99: Napster version 1.0
 - 5/99: company founded
 - 12/99: first lawsuits
 - 2000: 80 million users
- **Napster history: the fall**
 - Mid 2001: out of business due to lawsuits
 - Mid 2001: dozens of decentralized P2P alternatives
 - 2003: growth of pay services like iTunes



**Shawn Fanning,
Northeastern freshman**

Napster Directory Service

- Client contacts Napster (via TCP)



- Provides a list of music files it will share
- ... and Napster's central server updates the directory

- Client searches on a title or performer

- Napster identifies online clients with the file
- ... and provides their IP addresses

- Client requests the file from the chosen supplier

- Supplier transmits the file to the client

- Both client and supplier report status to Napster

Napster Properties

- **Server's directory continually updated**
 - Always know what music is currently available
 - Point of vulnerability for legal action
- **Peer-to-peer file transfer**
 - No load on the server
 - Plausible deniability for legal action (but not enough)
- **Bandwidth**
 - Suppliers ranked by apparent bandwidth and response time

Napster: Limitations of Directory


- File transfer is decentralized, but locating content is highly centralized
 - Single point of failure
 - Performance bottleneck
 - Copyright infringement
- So, later P2P systems were more distributed
 - Gnutella went to the other extreme...

Peer-to-Peer Networks: Gnutella


- Gnutella history
 - 2000: J. Frankel & T. Pepper released Gnutella
 - Soon after: many other clients (e.g., Morpheus, Limewire, Bearshare)
 - 2001: protocol enhancements, e.g., “ultrapeers”
- Query flooding
 - Join: contact a few nodes to become neighbors
 - Publish: no need!
 - Search: ask neighbors, who ask their neighbors
 - Fetch: get file directly from another node




Gnutella: Search by Flooding



Gnutella: Search by Flooding



Gnutella: Search by Flooding



Gnutella: Pros and Cons

- **Advantages**
 - Fully decentralized
 - Search cost distributed
 - Processing per node permits powerful search semantics
- **Disadvantages**
 - Search scope may be quite large
 - Search time may be quite long
 - High overhead, and nodes come and go often

Peer-to-Peer Networks: KaAzA

- KaZaA history


- 2001: created by Dutch company (Kazaa BV)
- Single network called FastTrack used by other clients as well
- Eventually protocol changed so others could no longer use it

- Super-node hierarchy

- Join: on start, the client contacts a super-node
- Publish: client sends list of files to its super-node
- Search: queries flooded among super-nodes
- Fetch: get file directly from one or more peers



“Ultra/super peers” in KaZaA and later Gnutella



KaZaA: Motivation for Super-Nodes

- Query consolidation
 - Many connected nodes may have only a few files
 - Propagating query to a sub-node may take more time than for the super-node to answer itself
- Stability
 - Super-node selection favors nodes with high up-time
 - How long you've been on is a good predictor of how long you'll be around in the future

Peer-to-Peer Networks: BitTorrent

- BitTorrent history
 - 2002: B. Cohen debuted BitTorrent
- Emphasis on efficient fetching, not searching
 - Distribute same file to many peers
 - Single publisher, many downloaders
- Preventing free-riding
 - Incentives for peers to contribute



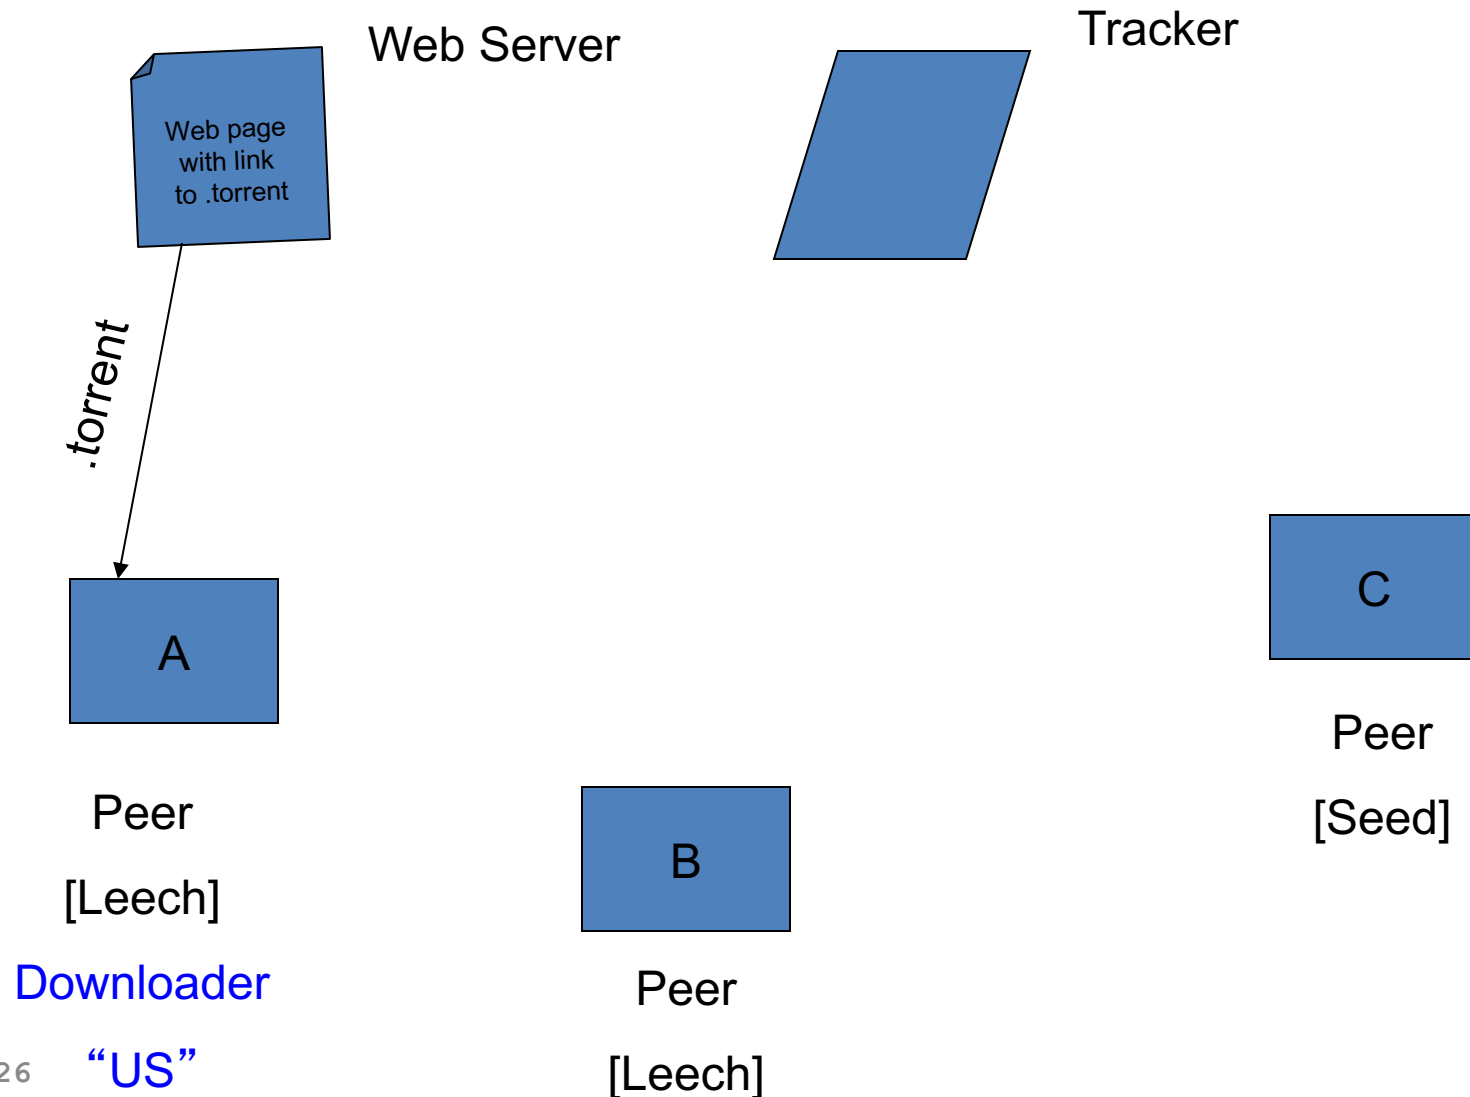
BitTorrent: Simultaneous Downloads

- Divide file into many chunks (e.g., 256 KB)
 - Replicate different chunks on different peers
 - Peers can trade chunks with other peers
 - Peer can (hopefully) assemble the entire file
- Allows simultaneous downloading
 - Retrieving different chunks from different peers
 - And uploading chunks to peers
 - Important for very large files

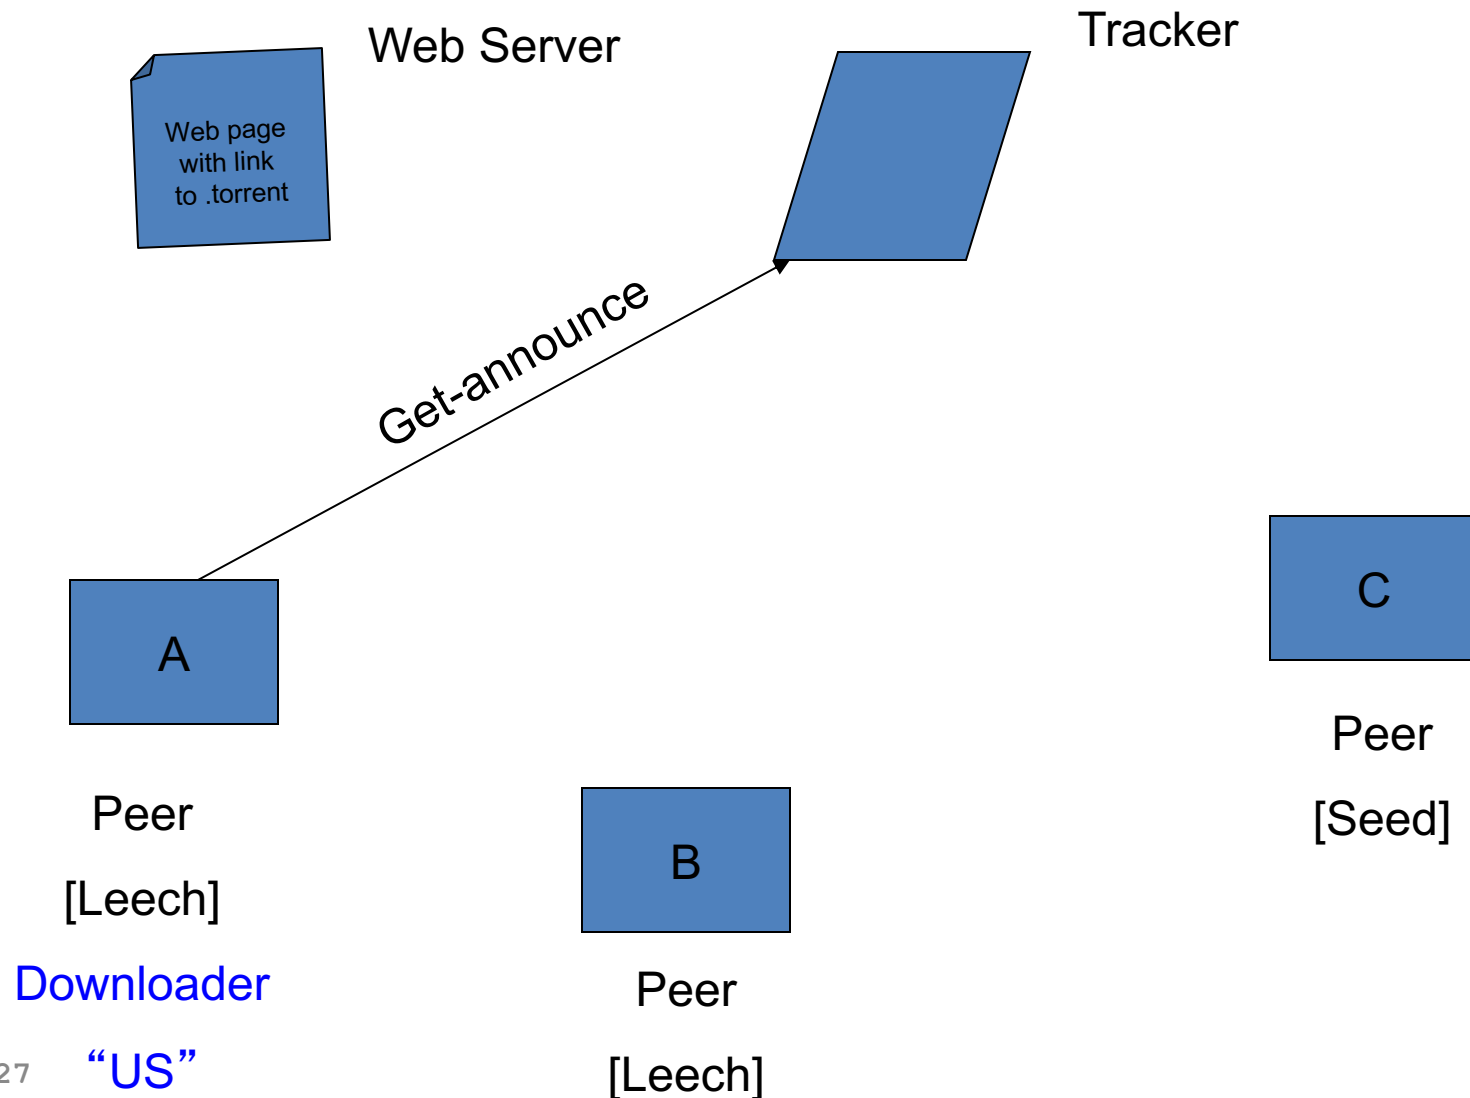
BitTorrent: Tracker

- Infrastructure node
 - Keeps track of peers participating in the torrent
 - Peers registers with the tracker when it arrives
- Tracker selects peers for downloading
 - Returns a random set of peer IP addresses
 - So the new peer knows who to contact for data
- Can have “trackerless” system
 - Using distributed hash tables (DHTs)

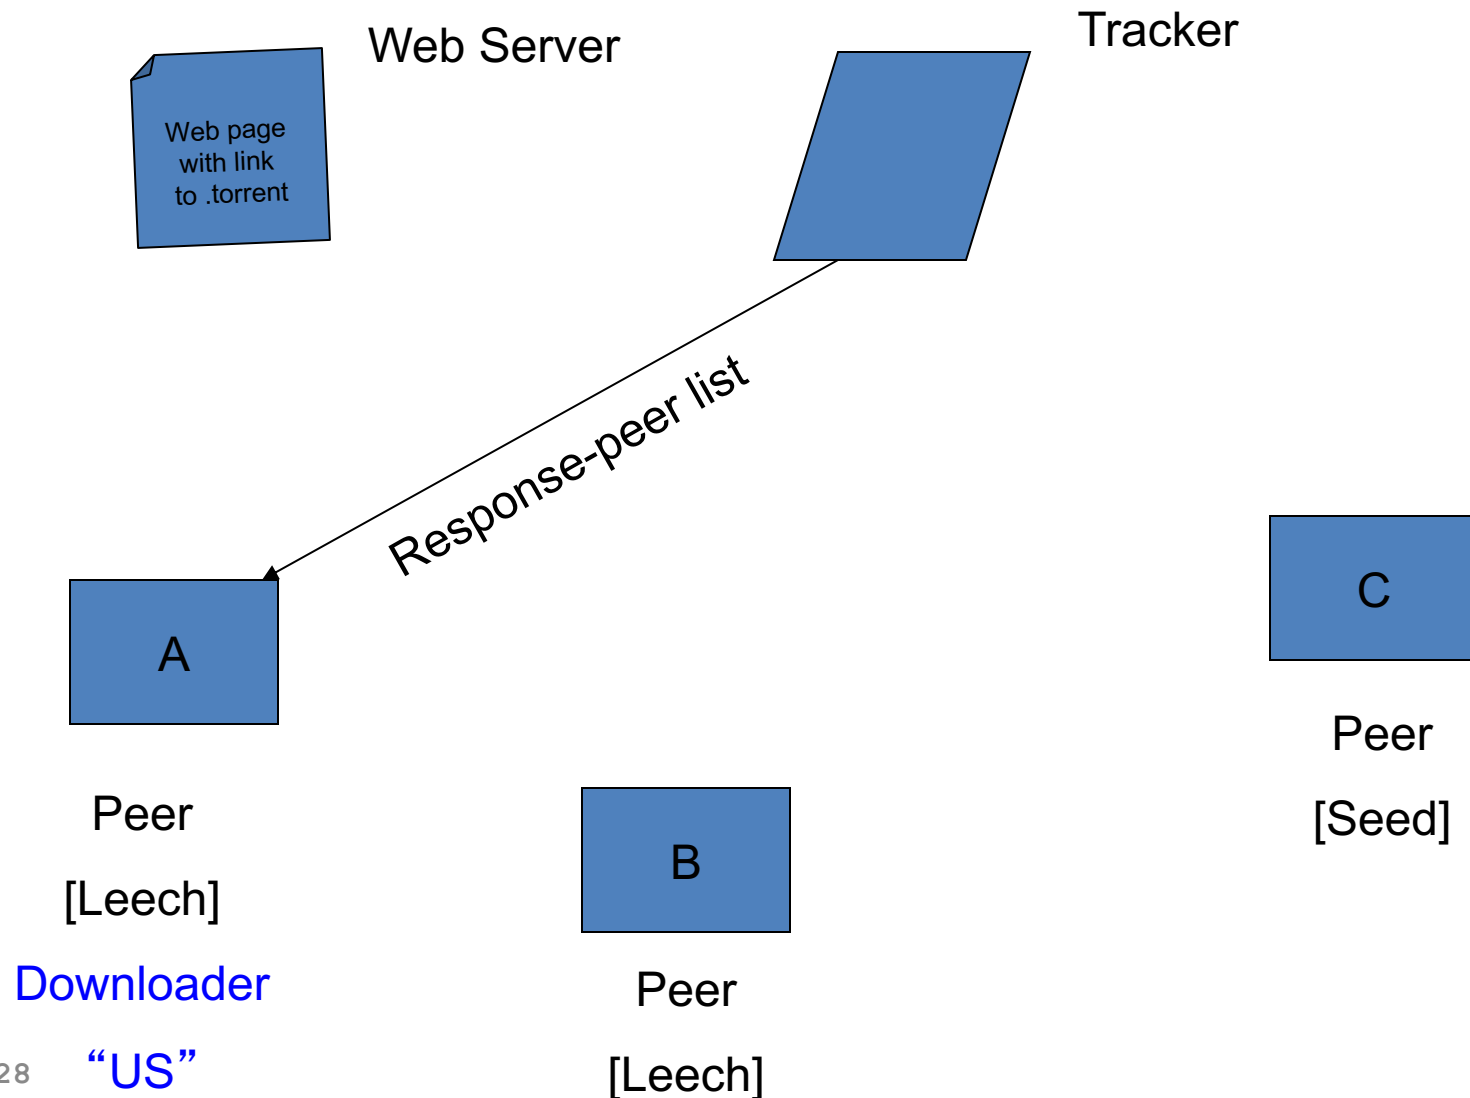
BitTorrent: Overall Architecture



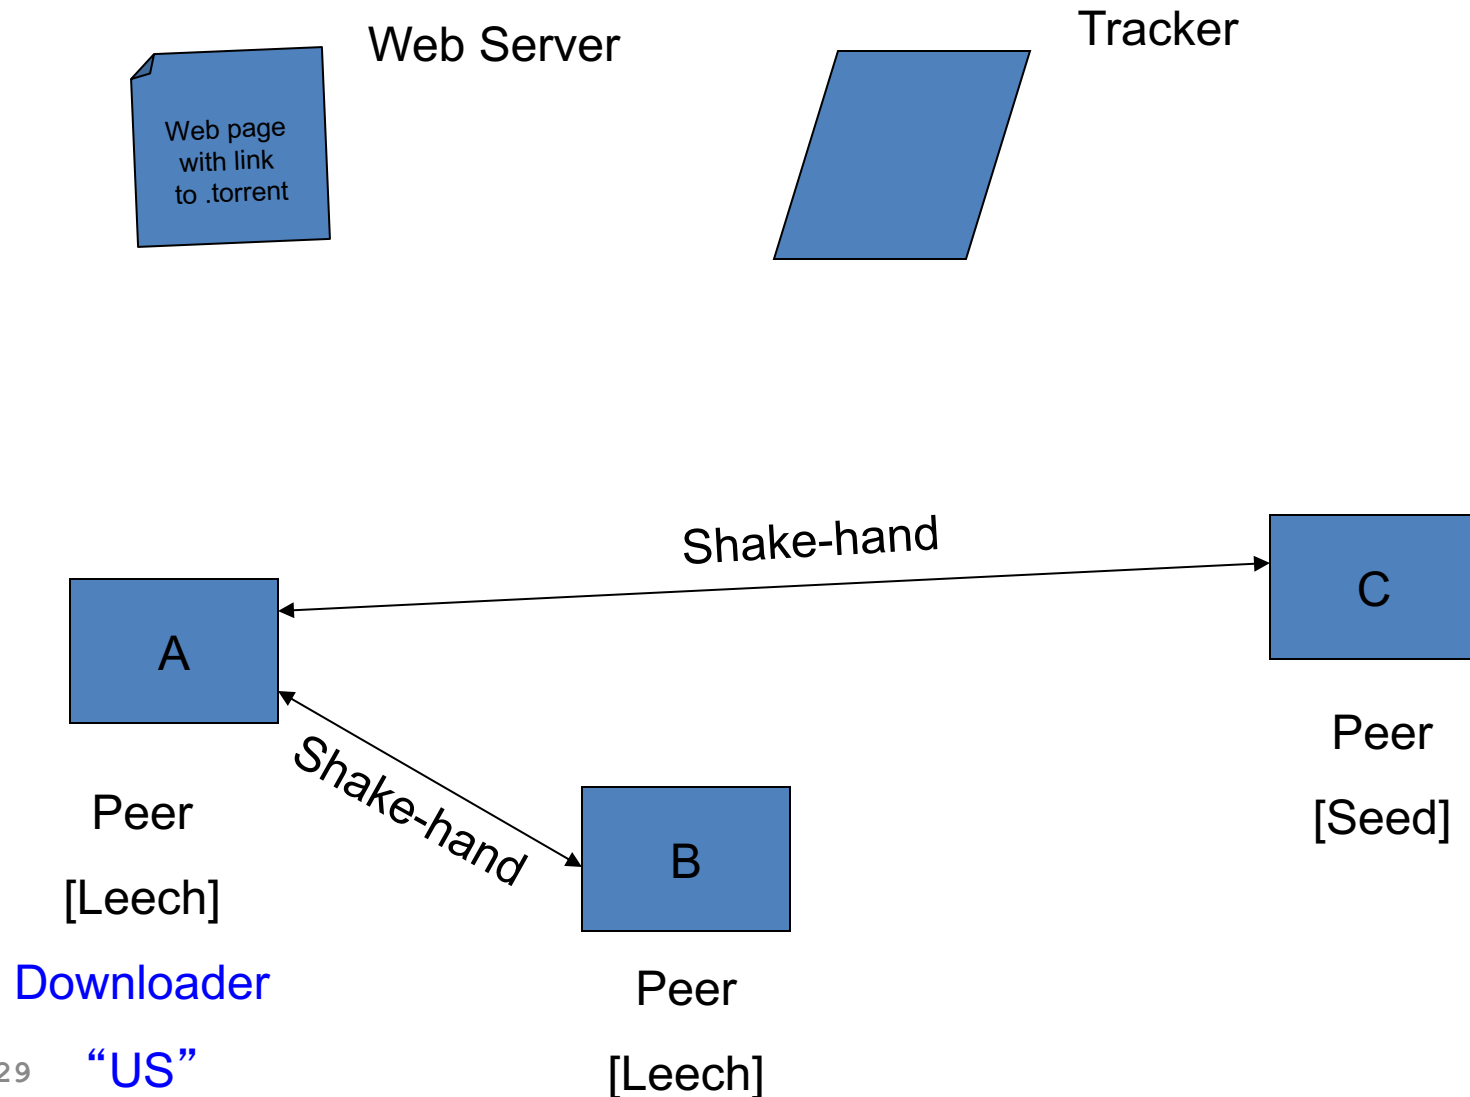
BitTorrent: Overall Architecture



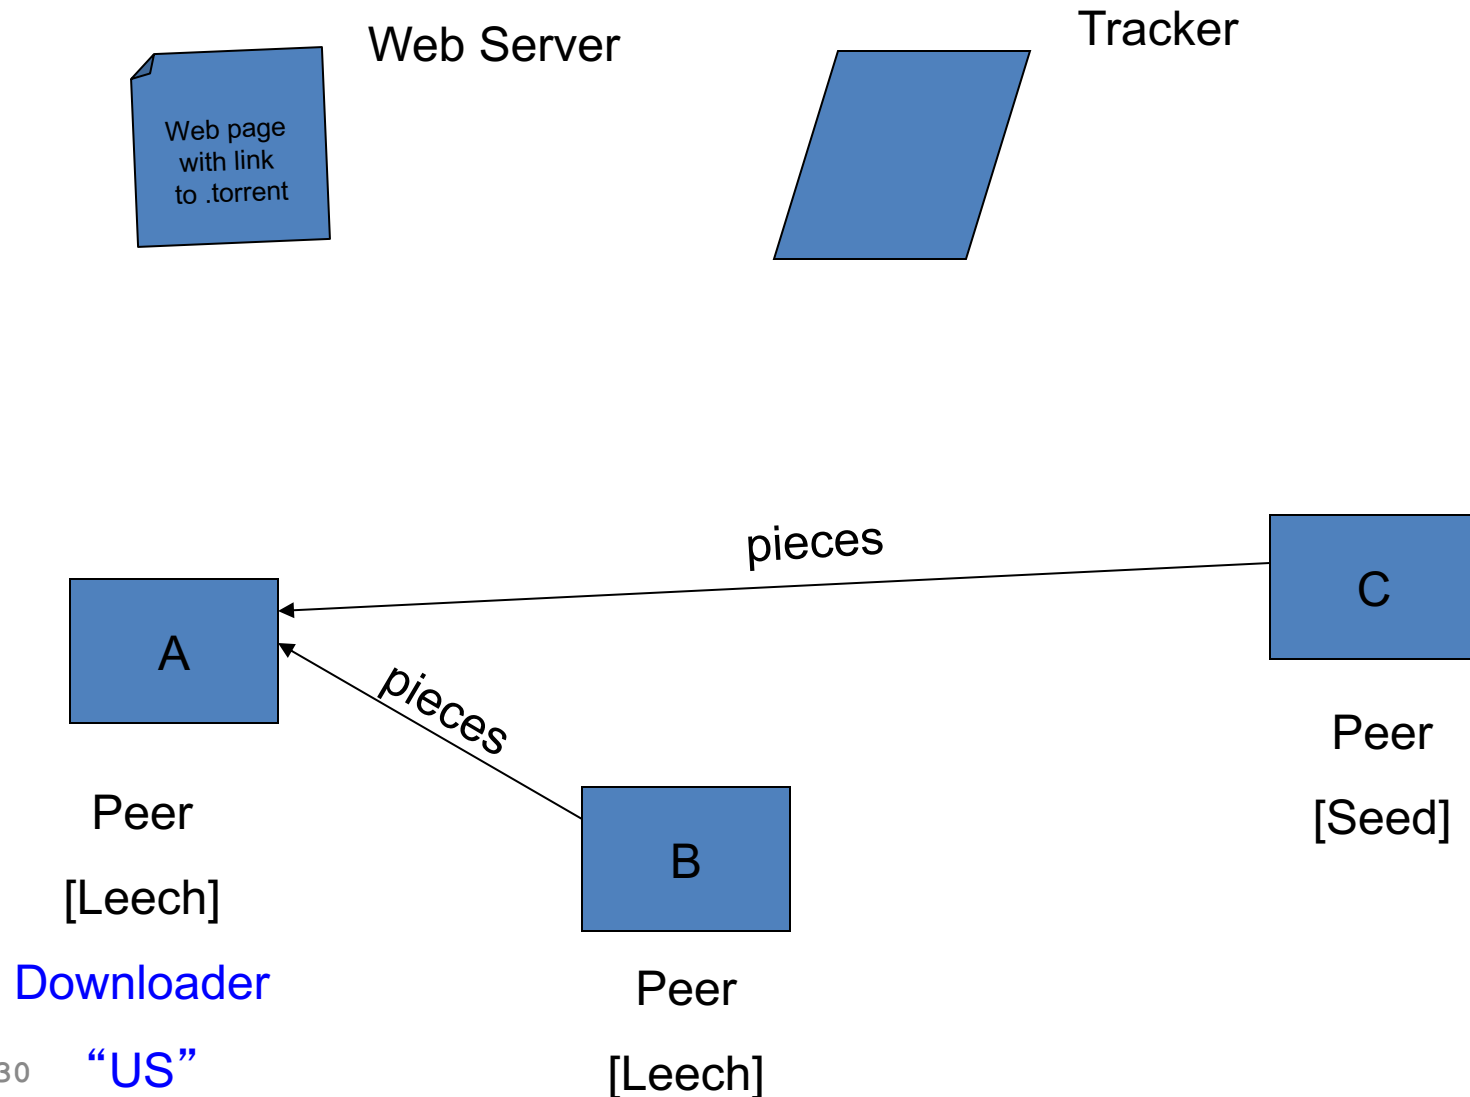
BitTorrent: Overall Architecture



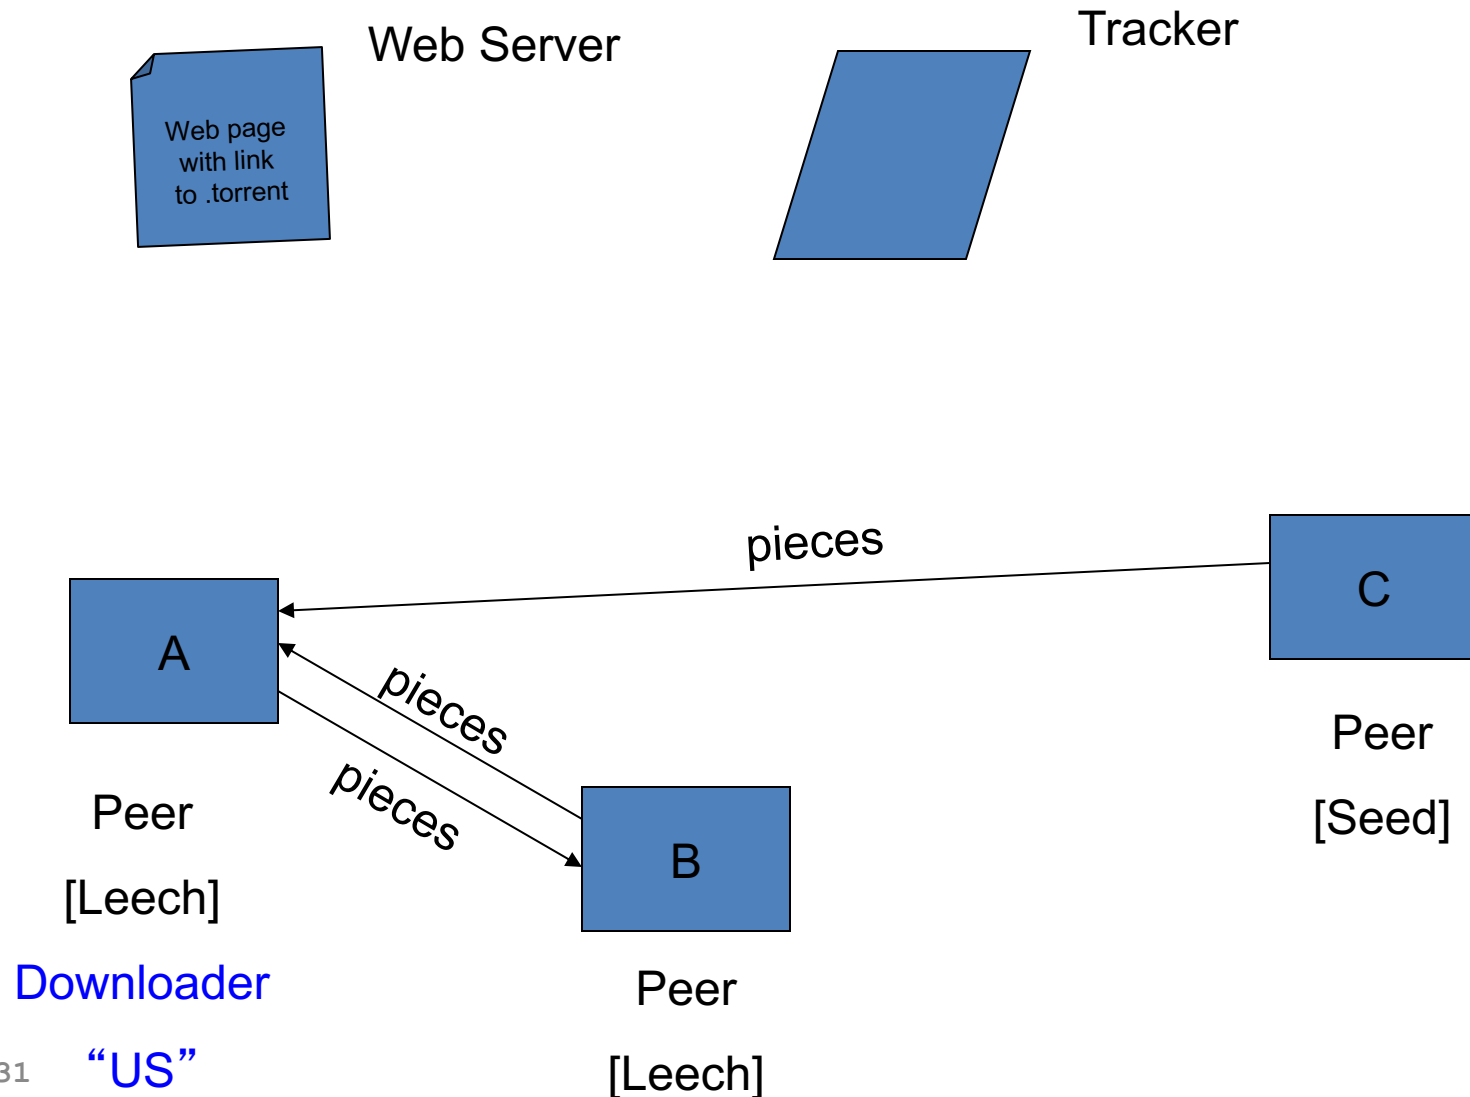
BitTorrent: Overall Architecture



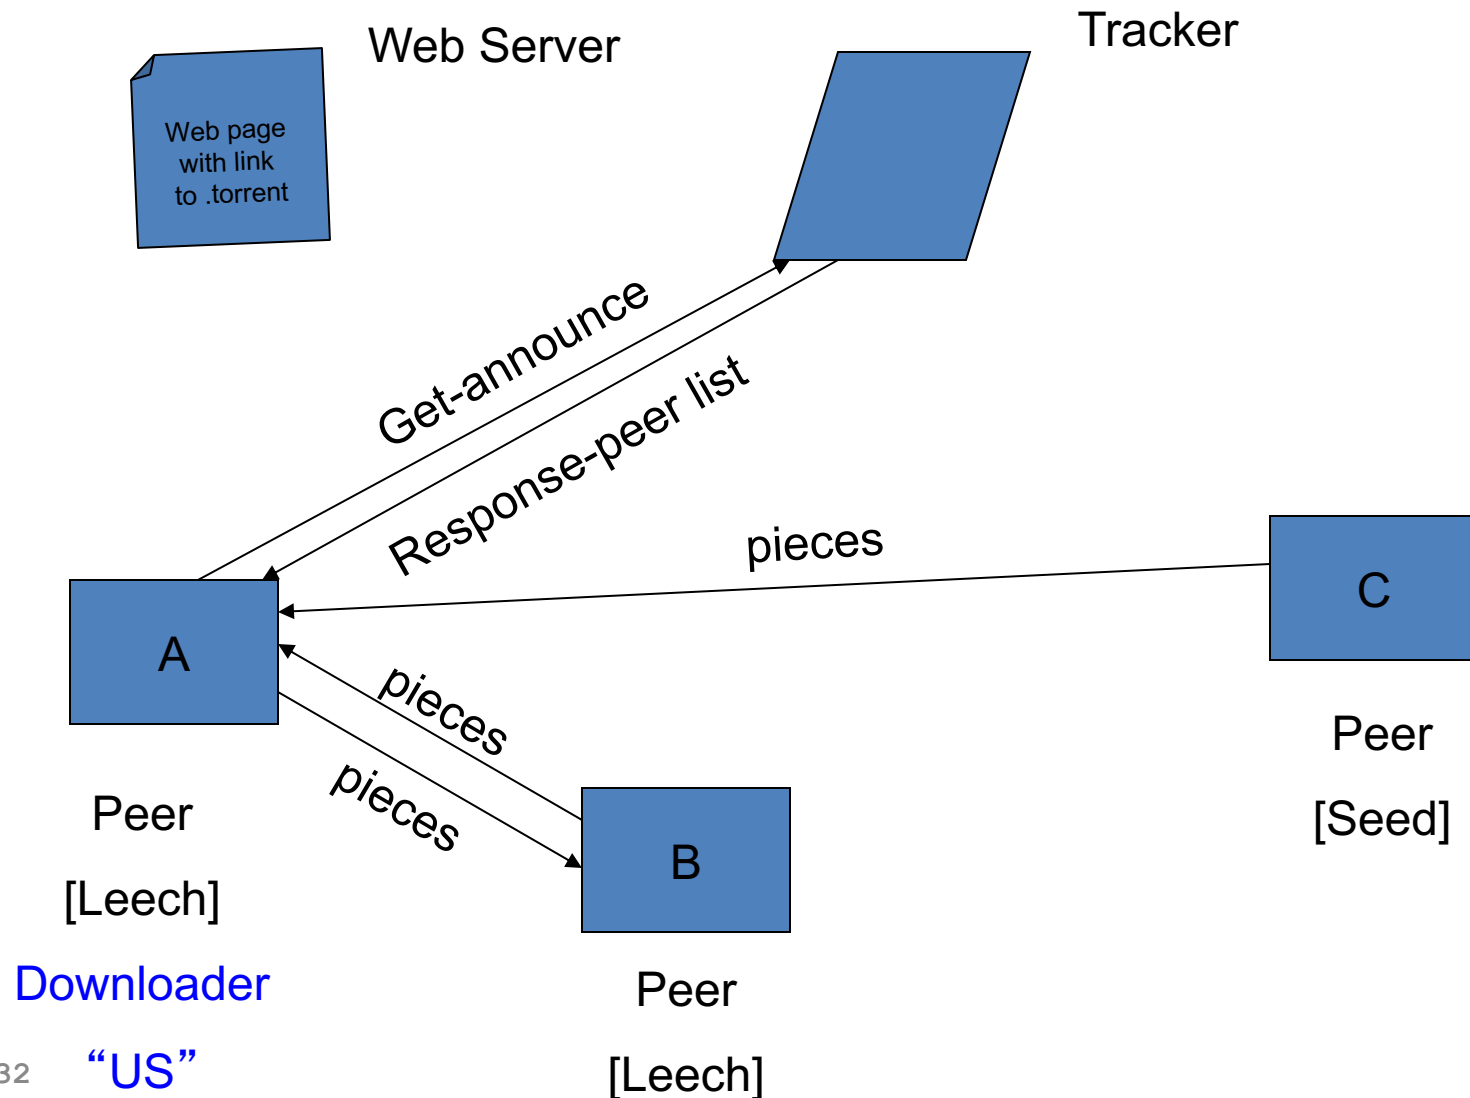
BitTorrent: Overall Architecture



BitTorrent: Overall Architecture



BitTorrent: Overall Architecture



BitTorrent: Chunk Request Order

- Which chunks to request?
 - Could download in order
 - Like an HTTP client does
- Problem: many peers have the early chunks
 - Peers have little to share with each other
 - Limiting the scalability of the system
- Problem: eventually nobody has rare chunks
 - E.g., the chunks need the end of the file
 - Limiting the ability to complete a download
- Solutions: random selection and rarest first

BitTorrent: Rarest Chunk First

- Which chunks to request first?
 - Chunk with fewest available copies (i.e., rarest chunk)
- Benefits to the peer
 - Avoid starvation when some peers depart
- Benefits to the system
 - Avoid starvation across all peers wanting a file
 - Balance load by equalizing # of copies of chunks

Free-Riding in P2P Networks

- Vast majority of users are free-riders
 - Most share no files and answer no queries
 - Others limit # of connections or upload speed
- A few “peers” essentially act as servers
 - A few individuals contributing to the public good
 - Making them hubs that basically act as a server
- BitTorrent prevent free riding
 - Allow the fastest peers to download from you
 - Occasionally let some free loaders download

Bit-Torrent: Preventing Free-Riding

- Peer has limited upload bandwidth
 - And must share it among multiple peers
 - Tit-for-tat: favor neighbors uploading at highest rate
- Rewarding the top four neighbors
 - Measure download bit rates from each neighbor
 - Reciprocate by sending to the top four peers
- Optimistic unchoking
 - Randomly try a new neighbor every 30 seconds
 - So new neighbor has a chance to be a better partner

BitTyrant: Gaming BitTorrent

- BitTorrent can be gamed, too
 - Peer uploads to top N peers at rate $1/N$
 - E.g., if $N=4$ and peers upload at 15, 12, 10, 9, 8, 3
 - ... peer uploading at rate 9 gets treated quite well
- Best to be the N^{th} peer in the list, rather than 1^{st}
 - Offer just a bit more bandwidth than low-rate peers
 - And you'll still be treated well by others
- BitTyrant software <http://bittyrant.cs.washington.edu/>
 - Uploads at higher rates to higher-bandwidth peers

Conclusions

- Finding the appropriate peers
 - Centralized directory (Napster)
 - Query flooding (Gnutella)
 - Super-nodes (KaZaA)
- BitTorrent
 - Distributed download of large files
 - Anti-free-riding techniques
- Great example of how change can happen so quickly in application-level protocols